

# **KAIST CS431:**

# **Concurrent Programming**

Instructor: Jeehoon Kang (<https://cp.kaist.ac.kr>)

# Logistics

- Homepage: <https://github.com/kaist-cp/cs431>
  - Read README.md carefully!
  - Announcement and question in the issue tracker (please watch the repo)
  - Office hours: Friday 9:15-10:15am
- Homework and attendance: <https://gg.kaist.ac.kr/16>
- Honor code: sign the KAIST CS honor code.
- Grading
  - Homework & project: 60%
  - Midterm & final exams: 40%
  - Attendance: ?

# Large Language Model Policy

- Large language models (LLMs): [ChatGPT](#), ...
  - We assume all of us can use ChatGPT 3.5.
- You can use LLMs for homework, study, ...
  - ChatGPT 4.0 doesn't help much for homework.
  - ChatGPT 3.5 helps much for studying CS431 materials.
- You cannot use LLMs for exams.
- We'll survey on your experience of LLMs for this course.

# Introduction




# Google DeepMind Challenge Match

8 - 15 March 2016







# The era of parallelism

- **Context: we need to do more and more computations**
  - Especially so in the era of AI/IoT
- **Trend: computers become more and more parallel**
  - Parallel resources: CPUs, memory, I/O, ...
- **Reason 1 (2005): the end of Dennard scaling**
  - Cannot increase the frequency of the circuits
  - Breakthrough: multi-core systems
- **Reason 2 (2018): the end of Moore's law**
  - Cannot increase # of transistors for fixed area
  - Breakthrough: accelerators (specialized, extremely parallel H/W)
- **How to coordinate parallel resources to achieve higher performance?**



# Parallel computing, theory and practice





# Concurrency: synchronizing parallel resources

- **Parallelism:** multiple resources
- **Concurrency: shared mutable resources (states)**
  - E.g. CPU, GPU, memory, server, database, datacenter, ...
  - Parallelism : concurrency = 찌넻뺑 : 꺑꺑소
    - Shared immutable resources: constant
    - Exclusive mutable resources: sequential
- **Example 1: lock-protected inode**
  - Inode: file system metadata
  - Serializes file accesses of multiple threads
- **Example 2: lock-free hash table**
  - Ensures correctness of simultaneous reads/writes of multiple threads

# Challenge in concurrency: nondeterminism

- **Challenge: combinatorially explosive nondeterminism**
- **Source 1: interleaving**
  - E.g., “X=1 || X=2”: the end memory depends on the order of execution
- **Source 2: optimization by hardware/compiler**
  - E.g., a=b=0 is possible in modern architecture by reordering:  
X=1 || Y=1  
a=Y || b=X
- **We need to tame the nondeterminism**

# Approaches to taming nondeterminism

- **Enclosing nondeterminism within safe API**
  - Hiding too low-level nondeterminism (e.g., interleaving of instructions)
  - While exposing high-level one (e.g., interleaving of queue operations)
  - Most people need to understand API, not implementation
  - E.g. safe API of locks, conditional variables, concurrent data structures
- **Reasoning with synchronization patterns in implementation**
  - Someone needs to implement lock, condvars, data structures, ...
  - Use only well-studied synchronization patterns
- **This course: learning API and implementation of concurrency libraries**

# Two modes of concurrency: “easy” and “difficult”

- **“Easy” lock-based concurrency**
  - Locks, conditional variables, ...
  - Pros & cons: simplicity & low scalability (losing parallelism opportunities)
  - Applicability: covering the majority 🤔 of use cases (in terms of # of lines)
- **“Difficult” lock-free concurrency**
  - **Theory:** semantics and reasoning principles  
(characterizing the nondeterminism)
  - **Tools:** synchronization patterns  
(building block for lock-free concurrency)
  - **Practice:** API and implementation of lock-free data structures  
(e.g., stack, queue, list, hash table, radix tree, balanced tree)




# General advices for concurrent programming

- **“Easy” concurrency is the first to study**
  - Go for “difficult” concurrency only when it’s a bottleneck
  - “Premature optimization is the root of all evil”
- **“Difficult” concurrency is not that difficult once you understand theory**
  - In other words, understanding theory may be difficult...
  - The key is taming the right amount of nondeterminism  
(Too much => scalability problem, too less => correctness problem)
- **Not so much to code, too much to debug**
  - Use sanitizer, stress testing, assert, line-by-line debugging, ...
  - Mathematical and PL thinking helps a lot

# Lock-based concurrency

## Part 1: motivation for safe API

# Context: shared-memory concurrency



- **Thread:** agent of execution reading/writing to shared memory
- **Shared memory:** shared storage of data
- **Other kinds of concurrency:** among CPU, memory, GPU/FPGA, persistent memory, distributed nodes, ...

# Lock-based shared-memory concurrency

- **Definition:** at any moment, a location is accessed by a single agent
- **Pros & cons:** simple & possibly inefficient
- **Mechanism:** locks (at any moment, only one thread holds the lock)

- **Examples**

- $r1=X \quad \parallel \quad r2=X$   
 $X=r1+1 \quad \parallel \quad X=r2+1$

// Not always  $X=2$ : unfortunate interleaving produces  $X=1$

- **L.acquire()**  $\parallel$  **L.acquire()**  
 $r1=X \quad \parallel \quad r2=X$   
 $X=r1+1 \quad \parallel \quad X=r2+1$   
**L.release()**  $\parallel$  **L.release()**

// Now always  $X=2$ : lock prevents unfortunate interleaving



# Why are locks “easy”?

- Recall: concurrency’s challenge is nondeterminism
  - Thread interleaving
  - Instruction reordering
- Lock constrains thread interleaving to acquire/release points
  - No interleaving between acquire and release
- Lock removes instruction reordering
  - Thread A’s release strictly happens before Thread B’s acquire
  - Threads are executed **AS IF** they are the same thread.
- **Lock reduces nondeterminism to the minimum.**

# Lock-based concurrency's low-level API

- **Lock.acquire():** Blocks until acquiring the lock.
- **Lock.try\_acquire():** Returns whether a lock is acquired. Doesn't block.
- **Lock.release():** Releases the acquired lock.
  
- **Challenge:** the API is extremely error-prone.
  - **Relating lock and resource:** users should access X only when L is held.
  - **Matching acquire/release:** users should release only acquired locks.
  
- **Consequence:** the API incurs high cost.
  - **Attention:** programmers should always be concerned with the API.
  - **Potential bugs:** there are typically many bugs that remain.

# Lock-based concurrency's high-level API

- **We want an easy-to-use, always-safe high-level API.**
  - Acquire/release are automatically matched.
  - Lock and resource are explicitly related.
- **Benefits of high-level API:** low cost with less attention and bugs
  - **Less attention:** Programmers don't need to worry about API misuses.
  - **Less bugs:** It's highly unlikely that a program using the API is buggy.
- **Design of high-level API (from C++/RAII)**
  - **LockGuard:** automatically releasing a lock using an RAII type
    - Lock.acquire() returns "lock guard"
    - When a lock guard is destructed, the corresponding lock is released.
  - **Lock<T> (= (Lock, T)):** relating a lock and a resource with a new type

# Lock-based concurrency's safe(ish) API: lock guard

- Lock guard: holding a lock

- [https://en.cppreference.com/w/cpp/thread/lock\\_guard](https://en.cppreference.com/w/cpp/thread/lock_guard)

```
#include <thread>
#include <mutex>
#include <iostream>

int g_i = 0;
std::mutex g_i_mutex; // protects g_i

void safe_increment()
{
    const std::lock_guard<std::mutex> lock(g_i_mutex);
    ++g_i;

    std::cout << std::this_thread::get_id() << ": " << g_i << '\n';

    // g_i_mutex is automatically released when lock
    // goes out of scope
}

int main()
{
    std::cout << "main: " << g_i << '\n';

    std::thread t1(safe_increment);
    std::thread t2(safe_increment);

    t1.join();
    t2.join();

    std::cout << "main: " << g_i << '\n';
}
```



# Lock-based concurrency's safe(ish) API: locked data

- **Locked data: a pair of lock and data**
- **Benefit: API mandates the inner data is (quite) safely protected by a lock**

- **// type**

```
template<typename T> class Lock<T> { RawLock lock; T data; }
```

- **// acquire and create lock guard**

```
LockGuard<T> Lock<T>::lock(this) {  
    this->lock.acquire(); LockGuard { this }  
}
```

- **// dereference data from lock guard**

```
&T LockGuard<T>::operator->(this) { &this->0.data }
```

- **// release automatically when guard is dropped**

```
LockGuard<T>::~~LockGuard() { this->0.lock.release(); }
```

# Lock-based concurrency's safe(ish) API: **not safe!**

- ```
// data: Lock<int>  
auto data_guard = data.lock();  
auto data_ptr = (int *) &data_guard;  
  
...  
// data_guard is dropped, lock is released  
*data_ptr = 666; // UNSAFE!
```
- **Root cause:** data\_ptr should not outlive data\_guard
- **Error-prone:** happening in the production code, causing lots of troubles
- **Solution: Rust's type system** based on **ownership and lifetime**
  - <https://github.com/kaist-cp/cs431/blob/main/src/lock/api.rs>
  - Rust implementation of lock has [proven-safe](#) API
  - => Let's first study Rust and then resume studying concurrency

# Lock-based concurrency

## Part 2: foundation for safe API in Rust

<https://docs.google.com/presentation/d/1LbiQ1Z3FTjp1144GRwEj3EPNj-RspAthlsq3a0PCQHw/edit#slide=id.p>

# Rust: **safe** systems programming language

- **Motivation:** achieving **safety & control** at the same time
  - Safety: compiled programs don't go wrong
  - Control: language supports low-level features
  - Prior art: C/C++ (unsafe)
- **Best fit for this course:** ownership and lifetime captures the essence of concurrency
- **Reading assignments:**
  - Read [the book](#). Homework 1 is about [the book's final project](#).
  - Read [Rust by example](#).
- **Programming assignments will be in Rust**
  - Set up programming environment on the [provided server](#).

# Rust example 1: iteration invalidation (1/3)

```
fn main() {  
    let v = vec![1, 2, 3];  
  
    let p = &v[1];  
    v.push(4);  
    println!("v[1]: {}", *p);  
}
```

- <https://play.rust-lang.org/?version=stable&mode=debug&edition=2018&gist=08f5870c40f7afdfd7a2fab9d7815f9f>
- This code would be compiled in C++, but it may fail at runtime
  - "v.push(4)" may relocate "v", Invalidating "p"
- This code is not compiled in Rust
  - "cannot borrow `v` as mutable because it is also borrowed as immutable"

# Rust example 1: iteration invalidation (2/3)

```
fn main() {  
    let mut v = vec![1, 2, 3];  
  
    let p = &v[1];  
    v.push(4);  
    println!("v[1]: {}", *p);  
}
```

- “v”: the owner of the vector
- “p”: immutably borrowing it  
from “let p = ...” to “println!(...)”
- “v.push(4)”: mutably borrowing it for the line
- Compile fails because type checker detects shared mutable accesses (SMA) to the vector
  - “p” and “v.push”
- It’s precisely the reason it may go wrong
- **Q: how to detect SMA’s?**

# Rust example 1: iteration invalidation (3/3)

```
fn main() {  
    let v = vec![1, 2, 3];  
  
    let p = &v[1];  
    v.push(4);  
    println!("v[1]: {}", *p);  
}
```

- Calculate each owner/borrower's **"lifetime"**
  - "v": L1-L5
  - "p": L3-L5
  - "v.push": L4
- List up the pairs of **overlapping lifetimes**
  - "v" and "p" (L3-L5), "v" and "v.push" (L4)
  - "p" and "v.push" (L4)
- Remove pairs of borrowee & borrower
  - "p" and "v.push" (L4) remained
- Remove pairs of immutable borrows
- The remaining pairs are **regarded as SMA's**
- **Statically sound:** detecting all SMA's at compile time
- Incomplete: not all pairs are actually SMA's



# Ownership for analyzing shared mutable accesses

- **“Ownership”**: the ability (of an agent) to access and destroy a resource
  - **Exclusive**: if I own a resource, no one else can own it
  - **Borrowable**: *mutably* borrowed by single agent or *immutably* by multiple
  - **Fit for concurrency**: each thread is an agent
- **Enforcing discipline**: no shared mutable accesses to a resource (by default)
  - **Static**: ownership discipline is enforced by types
  - **Easy to use**: compilers will report every violation of the discipline
  - **Correct**: if type-checked, a program doesn't go wrong
- **Bending the discipline** with “*interior mutability*”
  - **Necessary**: shared mutable accesses are inevitable in concurrency
  - **Modular**: enveloping implementation within safe API  
(“as if” there are no shared mutable accesses)

## Rust example 2: RefCell (1/4)

- **Context:** it is unrealistic to forbid SMA's altogether for, e.g., concurrency
- **Solution:** interior mutability
  - Enveloping SMA's in a safe API **as if** there are no SMAs
- **Example:** RefCell<T>
  - Checking ownership at runtime (not compile time)
  - RefCell<T>::try\_borrow(), RefCell<T>::try\_borrow\_mut():  
trying to borrow the inner value, immutably or mutably (resp.)
  - <https://doc.rust-lang.org/stable/std/cell/struct.RefCell.html>  
<https://doc.rust-lang.org/book/ch15-05-interior-mutability.html>

## Rust example 2: RefCell (2/4)

```
fn f1() -> bool { true }
fn f2() -> bool { !f1() }

fn main() {
    let mut v1 = 42;
    let mut v2 = 666;

    let p1 = if f1() { &v1 } else { &v2 };

    if f2() {
        let p2 = &mut v1;
        *p2 = 37;
        println!("p2: {}", *p2);
    }

    println!("p1: {}", *p1);
}
```

- <https://play.rust-lang.org/?version=stable&mode=debug&edition=2018&gist=c07efb0ed16980ef85d09568382114f9>
- Suppose f1() and f2() are complex and yet exclusive conditions (not f1() && f2())
- Safe because p1 and p2 are not aliased
- Compile error because the type checker cannot deduce the safety due to the complexity of conditions
- “cannot borrow `v1` as mutable because it is also borrowed as immutable”

## Rust example 2: RefCell (3/4)

```
use std::cell::RefCell;
```

```
fn f1() -> bool { true }
```

```
fn f2() -> bool { !f1() }
```

```
fn main() {
```

```
    let v1 = RefCell::new(42);
```

```
    let v2 = RefCell::new(666);
```

```
    let p1 = if f1() { &v1 } else { &v2 }
```

```
        .try_borrow().unwrap();
```

```
    if f2() {
```

```
        let mut p2 = v1
```

```
            .try_borrow_mut().unwrap();
```

```
        *p2 = 37;
```

```
        println!("p2: {}", *p2);
```

```
    }
```

```
    println!("p1: {}", *p1);
```

```
}
```

- <https://play.rust-lang.org/?version=stable&mode=debug&edition=2018&gist=00e64c84fc7b31b5080c4d386add8e9e>
- Ownership is checked at runtime (try\_borrow(), try\_borrow\_mut())
- Compiled and executed as expected  
“p1: 42”
- If f1() && f2(), try\_borrow\_mut() fails at runtime (not compile time)
- “thread 'main' panicked at 'called `Result::unwrap()` on an `Err` value’”

## Rust example 2: RefCell (4/4)

- **Interior mutability:** encapsulating SMA's in a non-SMA type
- **Safe API:** virtually w/o SMA's
  - `pub fn try_borrow_mut(&self) -> Result<RefMut<T>, BorrowMutError>`  
(**immutably** borrowing self)
- **Potentially unsafe implementation:** w/ SMA's
  - `... unsafe { &mut *self.value.get() }, ...`  
(<https://doc.rust-lang.org/1.63.0/src/core/cell.rs.html#1732>)
  - “**Unsafe**”: bridge between API w/o SMA's and impl. w/ SMA's  
(needs manual inspection, should be explicitly annotated)

# Rust example 3: Lock

- <https://github.com/kaist-cp/cs431/blob/main/src/lock/api.rs#L105>

```
102 impl<'s, L: RawLock, T> Deref for LockGuard<'s, L, T> {  
103     type Target = T;  
104  
105     fn deref(&self) -> &Self::Target {  
106         unsafe { &*self.lock.data.get() }  
107     }  
108 }
```

- `// data: Lock<int>`  
`let data_guard = data.lock();`  
`let data_ref = data_guard.deref();`  
`...`  
`drop(data_guard); // lock is released`  
**`*data_ref = 666; // NOT COMPILED: deref target shall not outlive guard`**

# Summary of Rust's ownership type

- **Motivation:** achieving **safety & control** over shared mutable resources
- **Key ideas:**
  - **Discipline:** disallowing shared mutable accesses by default
  - **Interior mutability:** allowing them in a controlled way
- **Benefits:**
  - Statically analyzing the safety of shared mutable accesses (both for sequential and concurrent programs)
  - Explicitly marking those code that needs manual inspection
- **What we'll do:** understanding lock-based concurrency with safe API

# Lock-based concurrency

## Part 3: safe API



# Rust concurrency libraries (1/3)

- **Potentially-unsafe implementations are enveloped within safe API**
  - If libraries are correct, the users don't need to worry about safety at all
- **Rust std**
  - [Thread](#): agent of execution
    - Safety: 'static closure (not function pointer), typed join handle
  - [Scoped thread](#): restricting thread's lifetime within a scope
    - Motivation: safe sharing of non-'static data
    - Safety: thread should be joined before the scope ('s) ends
  - [Arc](#): reference counter, immutably sharing data among multiple threads
    - Safety: Deref, not DerefMut
  - [Send](#): transferable to other thread
    - **Implementers**: `usize`, `&usize`, `Arc<T>`, `&Arc<T>` (but not `Rc<T>`, `&Rc<T>`)
  - [Sync](#): concurrently accessible from multiple threads
    - **Implementers**: `usize`, `Arc<T>` (but not `Rc<T>`)
    - **Property**: ``T: Sync`` if and only if `&T: Send`

# Rust concurrency libraries (2/3)

- [CS431 Lock API](#): a safe API for lock (see below for implementations)
- **Lock<L: RawLock, T>**: owns T that is protected by an L lock
  - Guarantee: the T object is not concurrently accessed (not code region)
  - Examples: Lock<SpinLock, Vec<usize>>, Lock<ClhLock, &'t TLS>
  - Property: Send + Sync if T is Send (i.e., meaningful only if T is Send)
- **LockGuard<'s, L: RawLock, T>**: proves the lock is acquired
  - Guarantee: the lock is held, T is accessible w/ Deref/DerefMut
  - RAI: it releases the lock when dropped
  - Property: Send if T is Send, Sync if T is Sync (i.e., transparent accessor)
- The API's guarantees/safety are proven w.r.t. Rust's ownership type system (as opposed to C/C++)

# Rust concurrency libraries (3/3)

- **More std**
  - [Mutex](#): mutual exclusion w/ various strategies
  - [Condvar](#): conditional variable, waiting for an event (condition)
    - Safety: `Condvar::wait()` gets `&mut MutexGuard`, forbidding reuse of protected data
  - [RwLock](#): reader-writer lock, allowing multiple readers **OR** one writer
- **crossbeam**
  - [Channel](#): sending/receiving values among threads
  - [CachePadded](#): align with 128 bytes
    - Motivation: to defeat “[false sharing](#)”
- **rayon**
  - [into\\_par\\_iter](#): executing a function for each element in parallel
    - Motivation: parallelism made easy

# Lock-based concurrency

## Part 4: implementation

# Several lock implementations

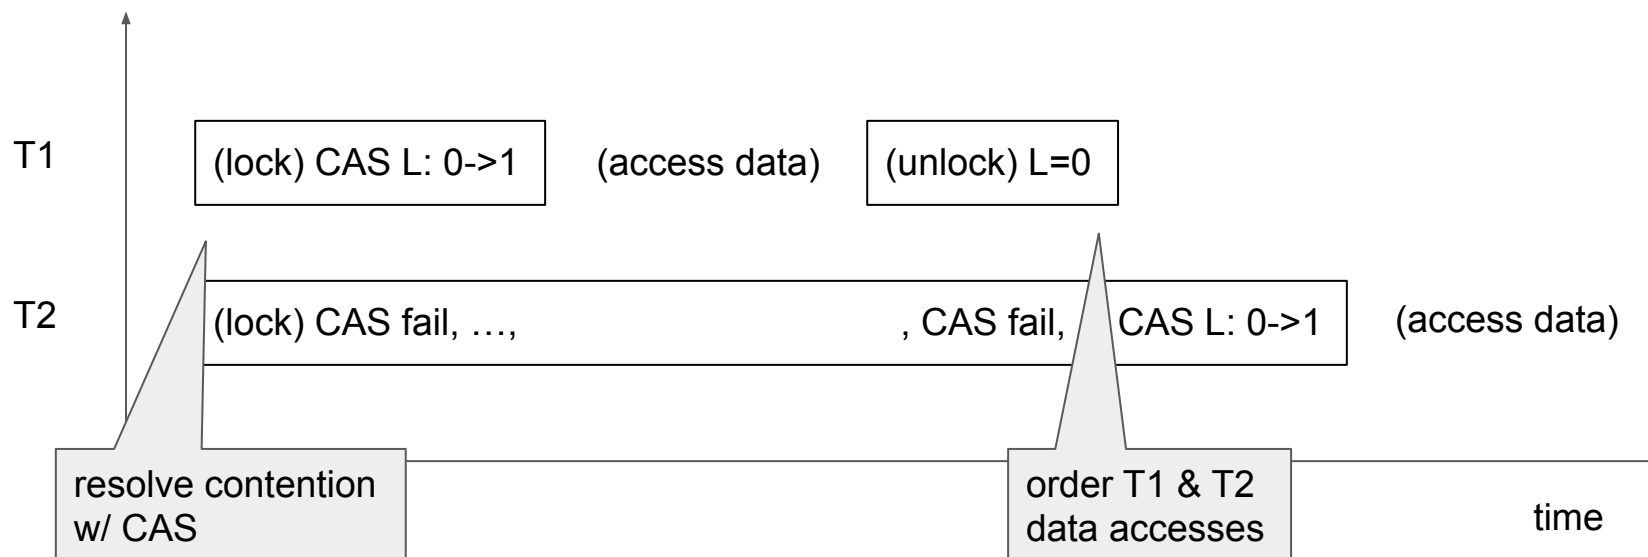
- **Implementations:** <https://github.com/kaist-cp/cs431/tree/main/src/lock>
  - Spinlock, ticket lock, CLH lock, MCS lock, MCS parking lock
- **RawLock trait:** defines the API of the “raw” lock
  - **APIs:** lock(), unlock(), Token, Default, Send, Sync
  - **Guarantee:** only one agent acquires the lock at a time
  - **Implementers:** spinlock, ticket lock, CLH lock, ...
- **Tradeoffs among locks**
  - Simple, fast (when uncontended), compact, scalable, fair, energy-efficient, ...

# Spinlock implementation

- <https://github.com/kaist-cp/cs431/blob/main/src/lock/spinlock.rs>  
(for now, ignore memory orderings: acquire, release, ...)
- ```
pub struct RawSpinLock {  
    inner: AtomicBool, // true means locked, false means unlocked  
}  
pub fn lock(&self) {  
    while self.inner.compare_and_swap(false, true).is_err() {} // rmw  
}  
pub fn unlock(&self) {  
    self.inner.store(false); // not rmw, thanks to exclusiveness of lock  
}
```

# Spinlock correctness

- ```
pub fn lock(&self) { while self.inner.cas(false, true, acquire).is_err() {} }  
pub fn unlock(&self) { self.inner.store(false, release); }
```
- If a lock has already been acquired, lock() will spin.
- Only one thread can hold a lock at a time (see below).



# The key ideas of the other locks

- <https://github.com/kaist-cp/cs431/tree/main/src/lock>
- **Guaranteeing mutual exclusion w/ CAS**
  - Ordering from the end of a critical section to the beginning of another
  - ``curr`` in ticket lock, **a new location** for each waiter in CLH/MCS lock
- **Guaranteeing fairness by ordering & waiting w/ different locations**
  - Ordering w/ fair instructions (e.g. swap, fetch-and-add)
  - Ticket lock: ordering w/ ``next`` and waiting w/ ``curr``
  - CLH/MCS lock: ordering w/ ``tail`` and waiting w/ **a new location**
- **Homework: reasoning the correctness of locks**



# Lock tradeoffs

- **Ticket lock:** guaranteeing fairness by ticket queueing
  - Lock order is decided by a fair instruction (fetch-add or swap) beforehand
  - Cons: a slightly complicated API (returning ticket)
- **CLH lock:** improving scalability by using per-critical section spinning location
  - Queue of spinning locations
  - Cons:  $O(n)$  space overhead, where  $n$  is the number of critical sections
- **MCS lock:** awaring NUMA by spinning on self-allocated location
  - Cons: possibly an additional compare-and-swap in unlock()
- **MCS parking lock:** reducing energy consumption by thread parking
  - Thread parking (intentionally blocking) instead of spinning
  - Cons: sacrificing performance in modest contention cases
- [A paper on performance evaluation](#)

# Lock questions

- **In ticket lock, why issuing a new ticket can be relaxed?**
  - In CLH lock, why swapping the tail can be relaxed?
  - In MCS (parking) lock, why swapping the tail can be relaxed?
- **In MCS parking lock, after unparked, why do we check if it's unlocked?**
- **In MCS parking lock, why is `thread` cloned?**
- **What else is in the literature?**
  - Reentrant lock: calling lock()s in a nested fashion
  - Reader-writer lock: allowing multiple readers or a single writer
  - Hierarchical lock: combining requests in a single NUMA node
  - Backoff strategy: spinning? yielding? parking? exponential backoff?
  - Conditional variable: guaranteeing order in addition to mutual exclusion
  - [https://docs.rs/parking\\_lot/0.11.0/parking\\_lot/](https://docs.rs/parking_lot/0.11.0/parking_lot/) (Rust impl of locks)
  - ...

# Lock-based concurrency


## Part 5: fine-grained locking

# Motivation: coarse-grained locking is unscalable

- **Coarse-grained locking:** protecting a large object w/ a single lock
  - Easy: the straightforward path towards concurrent programming
  - Unscalable: all accesses to the object are serialized w/o parallelism.
- **Fine-grained locking:** protecting many small objects w/ separate locks
  - More scalable: accesses are “distributed” to multiple locks  
(ideally, minimal synchronization overhead)
- **Drawbacks** of fine-grained locking (compared to coarse-grained one)
  - Single-thread overhead (ideally, only modest)
  - Complexity
- **Key design questions**
  - What should be done in sequence? -> lock required
  - What can be done in parallel? -> lock not required

# Example: lock-coupled linked list (structure)


- Each node has a lock
- When accessing the “next” pointer, lock should be held
  - A node’s lock protects the node’s next pointer
- Acquire the next node’s lock before releasing the current’s (“hand-over-hand”)
  - Make sure the current node is not detached



# Example: lock-coupled linked list (operations)


- **Insert B between A and C**

- Acquire A's lock
- Read A.next (=C)
- Allocate B with B.next=C
- Write A.next = B



- **Remove B between A and C**

- Acquire A's lock
- Read A.next (=B)
- Acquire B's lock
  - To ensure B.next is in the list
- Read B.next (=C)
- Write A.next=C
- (Release B's lock and) free B



# Lock-based concurrency

## Part 5: managing multiple locks w/ BoC

# Deadlock and livelock bugs

- **Deadlock:** each thread is blocked by another
  - e.g., A holding L1, acquiring L2  
B holding L2, acquiring L1
  - Countermeasure: detecting deadlock and killing an operation(transaction)
- **Livelock:** operations keeps getting killed without meaningful progress
  - A: L1 -> L2 -> L3  
B: L2 -> L3 -> L1  
C: L3 -> L1 -> L2
  - B holding L2, C holding L3/L1 -> C killed ->  
A holding L1, B holding L2/L3 -> B killed ->  
C holding L3, A holding L1/L2 -> A killed -> ...

**How to fight with them?**



# Deadlock and livelock avoidance


- Order locks: L1, L2, L3
  - e.g., according to their pointer address
- Acquire locks in the order
  - e.g., A holding L1, acquiring L2  
B holding L2, **acquiring L1 => NO!!!**
  - A: L1 -> L2 -> L3  
B: L2 -> L3 -> **L1 => NO!!!**  
C: **L3 -> L1 -> L2 => NO!!!**
- Theorem: no deadlock/livelock happens
  - Proof sketch: the thread holding the biggest lock will go unblocked

# BoC: Behavior-oriented Concurrency (Basics)

- <https://dl.acm.org/doi/10.1145/3622852>
- Motivations
  - Ordering locks automatically
  - Coupling lock-based synchronization and thread-based parallelization (cf: coupling synchronization and permission in high-level lock API)  
**=> Natural representation of task dependencies**
- Example

List. 5. Spawn a behaviour that requires both accounts

```
1  transfer(src: cown[Account], dst: cown[Account], amount: U64) {  
2    when (src, dst) { // withdraw and deposit  
3      if (src.balance >= amount && !src.frozen && !dst.frozen) {  
4        src.balance -= amount;  
5        dst.balance += amount;  
6      } } }
```

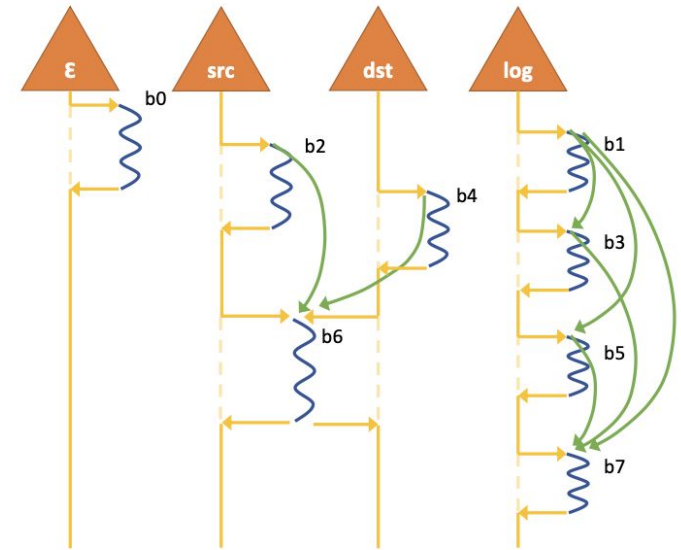


# BoC: Behavior-oriented Concurrency (Semantics)

- A behavior is **queued** for the (multiple) Cowns it uses
- Queueing guarantees inter-behavior dependencies and mutual exclusion

List. 8. Creating an accurate log

```
1  main(src: cown[Account], dst: cown[Account],
2      log: cown[OutStream]) { /* b0 */
3      when(log) { /* b1 */ log.log("begin") }
4      when(src) { /* b2 */ ...
5          when(log) { /* b3 */ log.log("deposit") }
6      }
7      when(dst) { /* b4 */ ...
8          when(log) { /* b5 */ log.log("freeze") }
9      }
10     when(src, dst) { /* b6 */ ...
11         when(log) { /* b7 */ log.log("transfer") }
12     } }
```



# BoC: Behavior-oriented Concurrency (Impl)

- Mutual exclusion for each Cown (similar to MCS lock's queueing)
- Each behavior requests Cowns in a fixed order (e.g., increasing order of addr)
- Each behavior acquires multiple cowns **atomically** by “2-phase enqueueing”
  - **Start phase** for all Cowns => **finish phase** for all Cowns
  - A behavior exclusively accesses a cown from start to finish phases
- Without 2-phase enqueueing, there can be a deadlock...
  - b1, b2 queued for c1, c2
  - c1 queue: b1, b2; c2 queue: b2, b1
  - => **DEADLOCK!**

# But what if we want even better performance?

- **Performance drawbacks of fine-grained locking**
  - Write operations have heavy synchronization costs due to lock ops
  - Read operations also writes to lock leading to cache invalidation
- **Alternative 1: advanced locks for read-mostly workloads**
  - reader-writer lock (“Rust concurrency libraries”),  
optimistic locking (not in semester), ... (still unscalable for writers)
- **Alternative 2: “lock-free” concurrent data structures**
  - Pros: even better performance
    - Write operations have lightweight synchronization costs
    - Read operations perform just reads w/o cache invalidation
  - Cons: *much* ( $\times 100$ ) more difficult than lock-based concurrency...

# **MIDTERM EXAM**

**Everything you've learned so far**

# Lock-free concurrency

## Part 1: definition

# Lock freedom is not about the existence of locks

- **What the hell?**
  - It was, but it isn't: the word's meaning has evolved over time.
- **Lock's limitations:** no guarantee of progress
  - Q: What if a thread holds a lock and sleeps forever?  
A: All threads waiting for the same lock stops making progress.
- **Lock freedom:** a progress guarantee
  - Definition: one of the ongoing operation is eventually completed
  - Intuition: the whole system's progress is guaranteed
  - Examples: Treiber stack, Michael-Scott queue, circular buffer, work-stealing deque, ...
- **Enemies:** thread stall due to I/O, crash, non-preemptive scheduling; locks (deadlock and livelock)



# Other progress guarantees

- **Obstruction freedom:** weaker than lock freedom
  - Definition: if only a single operation is run, it is eventually completed
  - Intuition: we can always “recover” an object to a stable state
- **Wait freedom:** stronger than lock freedom
  - Definition: every ongoing operation is eventually completed
  - Intuition: every thread’s progress is guaranteed
  - Examples: [Yang-Mellor-Crummey queue](#), ...
- $\text{Wait-free} \subseteq \text{lock-free} \subseteq \text{obstruction-free} \subseteq \text{“nonblocking”}$
- Reference: [Herlihy and Shavit. On the Nature of Progress. OPODIS'11](#)

# Key idea for lock freedom: single-instruction commit

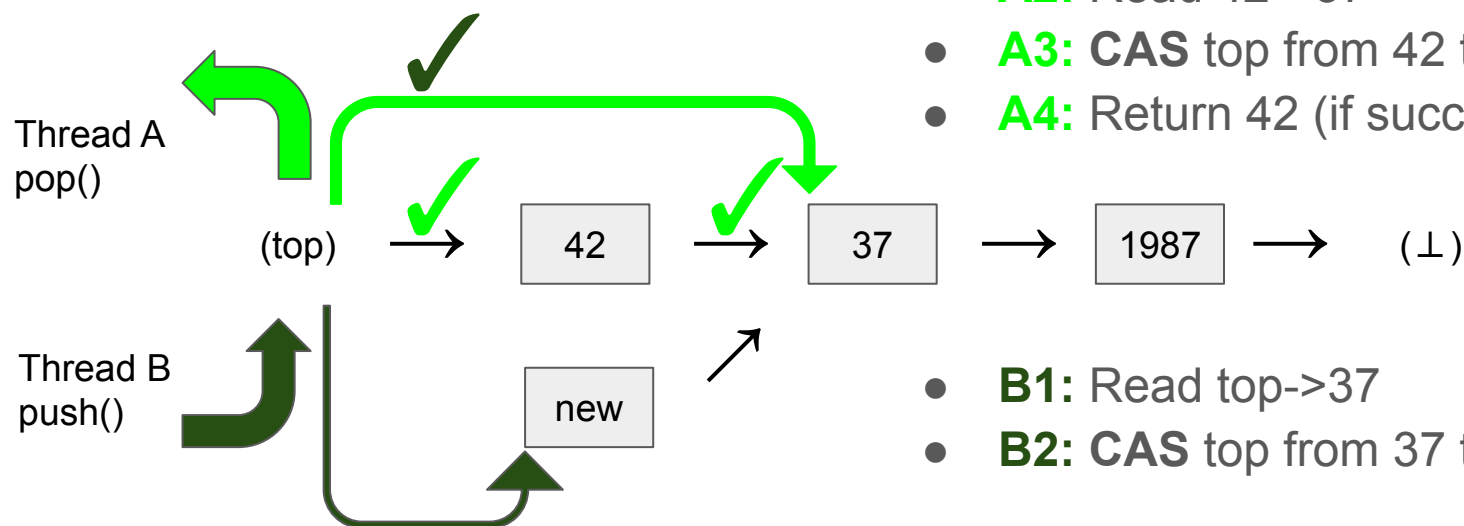
- **Leveraging architectural guarantee of lock freedom**
  - At any moment, at least one CPU core executes its instruction.
- **Designating an RMW instruction as an operation's commit point**
  - Definition: atomically reading from & writing to a location (e.g., CAS)
  - Property: capable of expressing synchronization protocols (e.g., spinlock)
  - History: single-instruction commit is RMW's motivation!
- **Defeating the “enemies” of lock freedom**
  - Thread pause/stall doesn't bother with a single RMW instruction.
  - Deadlock/livelock doesn't happen due to the absence of locks.
- **Side benefit: scalability** (reducing contention to a single instruction)

# Lock-free concurrency

## Part 2: data structures

# Example: Treiber's stack

- Singly linked list w/ list head = stack top
- <https://github.com/kaist-cp/cs431/blob/main/src/lockfree/stack.rs>




Push/pop synchronizes with CAS on the top location

# Questions for Treiber's stack

- Why data: `ManuallyDrop<T>` in `Node<T>`?
  - Because data's ownership is returned by `pop()`,  
while `Node<T>` will be deallocated later
  - Common in lock-free data structures

# Example: Michael-Scott's queue

- **Singly linked list** w/ list head/tail = queue pop/push ends (1 dummy at head)
- **Tail pointer:** push() doesn't need to traverse from the beginning
- <https://github.com/kaist-cp/cs431/blob/main/src/lockfree/queue.rs>
- **Pop(): CAS** on head (the same with Treiber's stack pop())
- **Push(): CAS** on tail from null to new



- **Challenge:** tail can be stale (e.g. right after pushing)
- **Solution:** relaxed invariant for tail (tail is reachable from head)

# Algorithm of Michael-Scott's queue

- **Push 1:** find the actual tail
  - Update the “tail” pointer if necessary
- **Push 2:** try to append a new node
  - If it fails, retry from the beginning
- **Push 3:** update the “tail” pointer
  - Perform a CAS from the old to the new tail
  - It's okay to fail the CAS (we just want tail pointer is sufficiently recent)
- **Pop 1: update the “tail” pointer if it points to the dummy node**
  - Ensure the head does not catch up to the tail
  - It's okay to fail the CAS (we just want tail pointer is sufficiently recent)
- **Pop 2: update the “head” pointer**
  - Perform CAS from the old to the new head
  - If it fails, retry from the beginning

# Questions for Michael-Scott's queue

- Why can a tail pointer be stale?
  - Right after a new node is inserted, the tail pointer becomes stale.
- Why have a dummy node at the head?
  - To make sure that the tail points to some node even for empty queues.
- In pop, why update the tail if it points to the dummy?
  - To make sure head doesn't "overtake" tail (why? see the GC part).
- Why are head & tail: `CachePadded<Node<T>>`?
  - To make sure that head and tail are not falsely shared.



# Example: sorted singly linked lists

- Linked list whose nodes (consisting of key, value, next) are **sorted by keys**
- <https://github.com/kaist-cp/cs431/blob/main/src/lockfree/list.rs>
- **Finding a node:** iterating list and finding the matching node
  - Used in lookup, insert, delete
- **Deleting a node:**
  - Step 1: tagging its next pointer by 0x1 (logical deletion)
  - Step 2: detaching it in another iteration (physical deletion)
- **Iteration strategies:** how to deal with logically deleted nodes?
  - Harris's: detaching consecutive logically deleted nodes at once
  - Harris-Michael's: detaching logically deleted nodes individually
  - Harris-Herlihy-Shavit's (wait-free): ignore logically deleted nodes
    - Can be used only for lookup (not insert/delete)

# Questions for sorted singly linked lists


- Why do logical deletion and then physical deletion?
  - To synchronize with insert right after the deleted node.
- What's the motivation of Harris-Michael's list?
  - To support hazard pointers; more on it later.
- What's the motivation of Harris-Herlihy-Shavit's list?
  - For wait-free lookup.
- Why can't HHS be used for insert/delete?
  - Cursor's "prev" may have already been deleted and cannot be updated.

# More examples

- **Circular buffer**
  - <https://people.mpi-sws.org/~dreyer/papers/gps/paper.pdf>
- **Hash table, binary trees (AVL/red-black), radix trees, ...**
- **Chase-Lev work-stealing deque**
  - <https://www.dre.vanderbilt.edu/~schmidt/PDF/work-stealing-dequeue.pdf>
  - <https://fzn.fr/readings/ppopp13.pdf>
  - <https://github.com/crossbeam-rs/crossbeam/tree/master/crossbeam-deque>
  - <https://github.com/jeehoonkang/crossbeam-rfcs/blob/deque-proof/text/2018-01-07-deque-proof.md>

# More about work stealing

- Originated from Cilk: <http://supertech.csail.mit.edu/cilk/>
- Dynamic balancing of workloads across parallel resources
  - By “stealing” work from the other threads
- (figures from [Lin Clark's blog post](#))



# Lock-based concurrency

## Part 5: optimistic locking

# Optimistic concurrency control (OCC)

- Observation: frequently, multiple operations are not conflicting with each other
- Idea: optimistically assumes the success of operation, posthumously recovers
- Primitive: sequence lock
  - <https://github.com/kaist-cp/cs431/blob/main/src/lock/seqlock.rs>
  - Observation: all read operations are not conflicting with each other
  - Idea: optimistically reads, posthumously validates with sequence number

# Sequence lock

- Optimistic reader-writer lock
- Writer: almost the same with spinlock
  - Managing sequence number (usize) instead of lock flag (boolean)
  - Even: consistent state between c.s., odd: inconsistent state inside c.s.
  - E.g. W(0): acquires 0->1, releases 2,  
W(2): acquires 2->3, releases 4, ...
- Reader: trying to read a consistent state (e.g. 2)
  - **Reading** sequence number @ beginning & end (should be even & same)
  - **Req 1:** W(0)'s end happens before R(2)'s beginning
  - **Req 2:** R(2) doesn't see W(2)'s writes as far as R(2) is validated

# Example: coarse-grained optimistic list

- algorithm (그림): global read lock -> find -> finish/upgrade
- memory reclamation problem
  - 위 그림 또는 GC 슬라이드 이용
  - RCU API만
  - old slide

[https://docs.google.com/presentation/d/1NMq08N1LUNDPuMxNZ-UMbdH13p8LXgMM3esbWRMowhU/edit#slide=id.ga54eefc4bc\\_1\\_651](https://docs.google.com/presentation/d/1NMq08N1LUNDPuMxNZ-UMbdH13p8LXgMM3esbWRMowhU/edit#slide=id.ga54eefc4bc_1_651)

- [https://git.fearless.systems/kaist-cp-class/cs431-private/-/blob/d92dac01a8036c7abdb6ccf13a90fcc338332c47/src/list\\_set/optimistic.rs](https://git.fearless.systems/kaist-cp-class/cs431-private/-/blob/d92dac01a8036c7abdb6ccf13a90fcc338332c47/src/list_set/optimistic.rs) (TODO: publish)
  - crossbeam\_epoch API
    - using Atomic<T> to explicitly allow race
    - reclamation



# Lock-free concurrency

## Part 3: garbage collection

In [a separate slide](#)

# Lock-free concurrency

## Part 4: linearizability

# Key objectives of concurrent data structures (CDS)

- **Progress:** guaranteeing the completion (or progress) of operations
  - Lock freedom: progress of at least one
  - Wait freedom: progress of everyone
- **Scalability:** showing better performance as the number of cores grow
  - Ideal: linear scaling
  - Reality: e.g. sublinear scaling after 16 threads
  - Key idea: reducing critical sections (fine-grained locks and CAS) & writes
- **Correctness:** working “as expected”
  - Safety: doesn’t go wrong (i.e., no segmentation fault)
  - Sequential specification: e.g. works like a queue
  - Synchronization: e.g. matched push and pop are synchronizing

# Scalability of CDS

- **Key idea 1: reduce contention by shrink lock protection scopes**
  - E.g. hand-over-hand locking, lock coupling, read-write locking
- **Key idea 2: reduce cache invalidations by avoiding writes**
  - E.g. “optimistic concurrency control”
  - E.g. avoiding writes in readers (especially for read-mostly workloads)
  - E.g. lightweight custom synchronization protocols
- **Case study: optimistic lock coupling**
  - Lock coupling + optimistic concurrency control
  - ~~Homework 2: Optimistic lock coupling for binary search tree~~

# Safety of CDS

- **Key idea: protect CDS w/ locks or more primitive synchronization**
  - Protecting a sequential DS with a global lock
  - Protecting a DS with more fine-grained locks
  - Protecting a DS with custom synchronization protocols
- **Specification: “linearizability”**  
(<https://github.com/jeehoonkang/crossbeam-rfcs/blob/deque-proof/text/2018-01-07-deque-proof.md>)

# Linearizability: the “right” correctness specification

- **Key idea:** the only complication of CDS over sequential DS is the order of operations (not the order of instructions)
- **Contextual refinement:** CDS works “as if” an abstract DS
  - Abstract DS: DS operations are single instructions
- **Linearizability:** a key lemma for contextual refinement
  - In an execution, there exists a total order  $R$  among CDS operations s.t.:
  - (VIEW) If  $o_1$  happens before  $o_2$ , then  $o_1 R o_2$ .
  - (SEQ) The results of operations are as if they are executed to sequential DS in the order  $R$ .
  - (SYN) E.g. a push operation happens before its matching pop operation.
  - More detail in [this document](#)

# How to prove linearizability?

- **Key question: how to find linearization order?**
  - Linearization point: a point in an operation that determines lin. order
  - Then what's the linearization point of each operation?
- **Key idea 1:** Write operation's commit point is its linearization point.
  - Recall: many lock-free DSs have single-instruction commit point
- **Key idea 2:** Read operation's *critical read* is its linearization point.
  - Example: [read of head = null] for a pop from the empty Treiber's stack
- **Challenges & research questions**
  - Some data structures don't have commit point (e.g. Chase-Lev deque).  
How to find the linearization order for such data structures?
  - How to *prove* linearization?

# Linearization examples

- **Treiber's stack**

- Key idea 1: DS operation order = the order of successful CAS
- Key idea 2: empty pop's order = between the surrounding pop and push

- **Michael-Scott's queue**

- Key idea 1: DS operation order = the order of successful CAS on the head & node's next pointers (but not on tail)

- **Harris-\*'s lists**

- Key idea 1: DS write operation order = the order of successful CAS/FAO  
On the head & node's next pointers
  - Except for detaching logically deleted nodes (it's not commit point)
- Key idea 2: DS read operation order = active research area
  - Cf. [paper1](#) [paper2](#)



# Towards formal verification of linearizability

- Video lecture by Dr. Derek Dreyer (MPI-SWS)
- RustBelt: Logical Foundations for the Future of Safe Systems Programming
- At Jane Street (2019)



# Relaxed-memory behaviors

## Part 1: semantics

# Motivation: lock-freedom reveals gory details

- Recall: concurrency is challenging due to nondeterminism arising from...
  - Interleaving: thread A and B's executions are interleaved; and
  - Reordering: thread A's instructions may be reordered.
- Reordering example  
DATA=42 || if FLAG==1  
FLAG=1 || assert(DATA==42) // may fail due to reordering!
- Lock would hide such reordering: X and Y are NOT concurrently accessed!
  - Reordering doesn't happen across lock/unlock.
  - Each variable (X or Y) is accessed only with holding a lock.
  - => all accesses to each variable is completely ordered
- Remaining questions
  - How to forbid reordering across lock/unlock?
  - How to forbid undesirable reordering in lock-free programming?

# Nondeterminism: the challenge of shared memory

- Memory: location -> byte w/ concurrent load/store instructions
- The single-most widely-used SMS
- **Challenge: highly nondeterministic**
- **Source 1 of nondeterminism: thread interleaving**
  - Load/store instructions of multiple threads are interleaved
  - Resulting in combinatorially explosive # of behaviors
- **Source 2 of nondeterminism: reordering**
  - Load/store instructions inside a single thread may be reordered
  - Resulting in unintuitive behaviors
- **Strategy: taming nondeterminism by forbidding unintended behaviors**

# Nondeterminism due to reordering

- **Motivation:** performing optimizations as if it's a sequential program
- **E.g.** message passing (FLAG is a shared location, e.g. AtomicUsize)  

```
DATA = 42;      ||  if FLAG.load() {  
FLAG.store(1); ||      assert(DATA == 42);  
                ||  }
```
- **Problem:** unintended behaviors due to reordering by compiler/hardware
  - Assertion failure due to reordering in the left thread or in the right thread
  - **Relaxed behaviors:** observable behaviors  
not captured in the “interleaving semantics”
- **Solution:** forbidding such a reordering with **ordering primitives**
  - **Fence:** fence(SC) between stores and between loads, or
  - **Access ordering:** FLAG.store(1, release) and FLAG.load(acquire)

# Enforcing Order in Spinlock (1/2)


- <https://github.com/kaist-cp/cs431/blob/main/src/lock/spinlock.rs>

- ```
22  ✓ unsafe impl RawLock for SpinLock {
23      type Token = ();
24
25  ✓  fn lock(&self) {
26      let backoff = Backoff::new();
27
28      while self
29          .inner
30          .compare_exchange(false, true, Acquire, Relaxed)
31          .is_err()
32      {
33          backoff.snooze();
34      }
35  }
36
37  unsafe fn unlock(&self, _token: ()) {
38      self.inner.store(false, Release);
39  }
40  }
```

- “Ordering::Acquire” and “Ordering::Release” enforces the order.

# Enforcing Order in Spinlock (1/2)

- // Thread 1  
DATA = 42  
LOCK = false (L.unlock()) // **RELEASE** to prevent reordering w/ above
- // Thread 2  
if (LOCK.cas(false, true)) // **ACQUIRE** to prevent reordering w/ below  
assert(DATA == 42)
- Release/acquire synchronization
  - If a release write is read by an acquire read,  
Then the write is **strictly happening before** the read.



# More relaxed behaviors due to reordering

- **Reordering:** unless accessing the same location,  
    **any** two load/store/rmw instructions can be reordered.
  - E.g. `X=1; Y=1 -> Y=1; X=1` (store-store reordering)
- **E.g. load hoisting** [`r1=r2=0` **allowed** by store-load reordering]
  - `X=1    ||    Y=1`  
    `r1=Y   ||   r2=X`
- **E.g. store hoisting** [`r1=r2=1` **allowed** by load-store reordering]
  - `r1=X   ||   r2=Y`  
    `Y=1    ||   X=1`
- **E.g. Java causality test cases** (#16, #19, #20 are wrong)  
    <http://www.cs.umd.edu/~pugh/java/memoryModel/CausalityTestCases.html>
- **Homework:** (informally) making sense of the above examples



# Forbidding reordering w/ ordering primitives

- **E.g. message passing w/ release/acquire synchronization**
  - DATA = 42;                      || if FLAG.load(**acquire**) {  
    FLAG.store(1, **release**); ||     assert(DATA == 42);  
                                  || }
    - **Release store**: forbidding reordering itself w/ earlier instructions
    - **Acquire load**: forbidding reordering itself w/ later instructions
- **E.g. message passing w/ sequentially consistent (SC) synchronization**
  - DATA = 42;                      || if FLAG.load(**relaxed**) {  
    **fence(SC)**;                    ||     **fence(SC)**;  
    FLAG.store(1, **relaxed**); ||     assert(DATA == 42); }
    - **SC fence**: forbidding any reordering across itself
    - **Relaxed**: imposing no orderings
- **Q: What is the precise meaning of relaxed behaviors and orderings?**

# Challenge: modeling relaxed behaviors & orderings

- **Approach 1: avoiding concurrent accesses altogether with locks**
  - DRF (data-race freedom): no concurrent accesses, no relaxed behaviors.
  - Problem: concurrent accesses are essential
- **Approach 2: constraining executions by “axioms”**
  - Representing an execution as a value-flow graph, validated w/ axioms
  - Problem 1: not operational semantics (less intuitive)
  - Problem 2: allowing the bad “out-of-thin-air” behaviors
- **Approach 3: modeling reorderings w/ operational semantics**
  - “Promising semantics”: <https://sf.snu.ac.kr/promise-concurrency/>
  - Operational semantics


# Promising semantics

- Kang et al. A promising semantics for relaxed-memory concurrency. POPL 2017
- Interleaving operational semantics modeling relaxed behaviors and orderings
- **Key idea 1:** modeling load hoisting w/ **multi-valued memory**
  - Allowing a thread to read an old value from a location
- **Key idea 2:** modeling read-modify-write w/ **message adjacency**
  - Forbidding multiple read-modify-writes of a single value
- **Key idea 3:** modeling coherence & ordering w/ **views**
  - Constraining a thread's behavior
- **Key idea 4:** modeling store hoisting w/ **promises**
  - Allowing a thread to speculatively write a value

# Promising semantics 1: multi-valued memory

- **Memory:** location  $\rightarrow$  list message, **message:** value & timestamp
- Threads may read an old value from a location (effectively hoisting loads)
- E.g. load hoisting [ $r1=r2=0$  **allowed** by reading old values from X and Y]

○ ●  $X=1$     ||    ●  $Y=1$   
○ ●  $r1=Y$  [0]    ||    ●  $r2=X$  [0]  
●                      ●




# Promising semantics 2: message adjacency

- **Message:** value & timestamp range
- Read-modify-writes should put the new msg. **adjacent** to the old one (no gap)
- A message occupies a range of timestamps (e.g. (10, 20])


- E.g. counter [r1=r2=0 **forbidden**]

○ ● r1=X.fetch\_add(1) 0 || ● r2=X.fetch\_add(1) 1




# Promising semantics 3: views (1/4)

- Multi-valued memory allows too many, unintended behaviors
- Needs to constrain the behaviors w.r.t. coherence and synchronization
- **View:** location  $\rightarrow$  timestamp (acknowledging messages for each location)  
representing acknowledgement of messages
  - **Per-thread view** for coherence
  - **Per-message view** for release/acquire synchronization
  - **A global view** for SC synchronization




## Promising semantics 3: views (2/4)

- **Per-thread view** representing a thread's acknowledgement of messages
- For modeling per-location coherence
  - RR coherence:  $X=1 \parallel r1=X; r2=X$  [ $r1=1, r2=0$  impossible]
  - RW coherence:  $r=X; X=1$  [ $r=0$ ]
  - WR coherence:  $X=1; r=X$  [ $r=1$ ]
  - WW coherence:  $X=1; X=2$  [ $X=2$  at the end]
- Reading/writing happens after the current thread's view
- Reading/writing changes the current thread's view



# Promising semantics 3: views (3/4)


- **Per-message view** representing the released view of the corresponding store
- For modeling release/acquire synchronization
- E.g. message passing (X=1 should be acknowledged after reading Y=1)
  - ● X = 1;                                   || ● if Y.load(acquire) {   1
  - Y.store(1, release);   || ●    assert(X == 1);   1
  - || ●    }





# Promising semantics 3: views (4/4)

- **A global view** representing the currently accumulated view of SC fences
- For modeling SC-fence synchronization (strict order among SC fences)
- After an SC fence, SC view and thread's view become the maximum of them
- E.g. message passing (X=1 should be acknowledged after reading Y=1)
  - ● X = 1;                      || ● if Y.load(relaxed) { 1
  - fence(SC);                 || ● fence(SC);
  - Y.store(1, relaxed); || ● assert(X == 1); 1
  -



# Promising semantics 4: promises (1/5)


- **Challenge: modeling store hoisting**  
(“A major open problem for PL semantics”, Batty et al., ESOP 2015)
- **Store hoisting w/o dependency** [ $r1=r2=1$  **allowed** by reordering in the right]  
     $r1=X \quad || \quad r2=Y$   
     $Y=r1 \quad || \quad X=1$
- **Store hoisting w/ dependency** [ $r1=r2=1$  **disallowed**, “out of thin air” (OOTA)]  
     $r1=X \quad || \quad r2=Y$   
     $Y=r1 \quad || \quad X=r2$
- **Store hoisting w/ syntactic dependency** [ $r1=r2=1$  **allowed** by compiler opt.]  
     $r1=X \quad || \quad r2=Y$   
     $Y=r1 \quad || \quad \text{if } r2==1 \{ X=r2 \} \quad // \text{ “if” should be taken for the behavior,}$   
                     $\text{else } \{ X=1 \} \quad // \text{ but looks like OOTA}$ 
  - compiler may (1) forward  $r2=1$  in the then branch and (2) hoist  $X=1$  out

## Promising semantics 4: promises (2/5)

- **Goal:** allowing the hoisting of **semantically independent writes** only
- We **will** cover promises in the final exam.
- **Idea:** “semantically independent writes” are **always writable** in the future
- **Mechanism**
  - A thread may speculatively write a value (“promise to write”)
  - A thread should **always be able to write its promises** in the future
- **Examples:** store hoistings w/ & w/o (syntactic) dependency


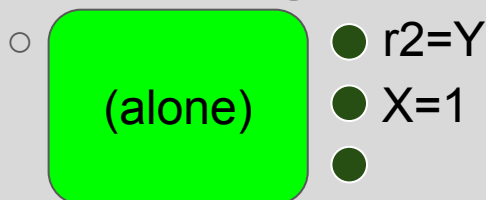
# Promising semantics 4: promises (3/5)

- **Store hoisting w/o dependency** [ $r1=r2=1$  **allowed** by reordering in the right]
  - ●  $r1=X$  || ●  $r2=Y$   
 $Y=r1$  ||  $X=1$



# Promising semantics 4: promises (3/5), certification

- Store hoisting w/o dependency [ $r1=r2=1$  **allowed** by reordering in the right]



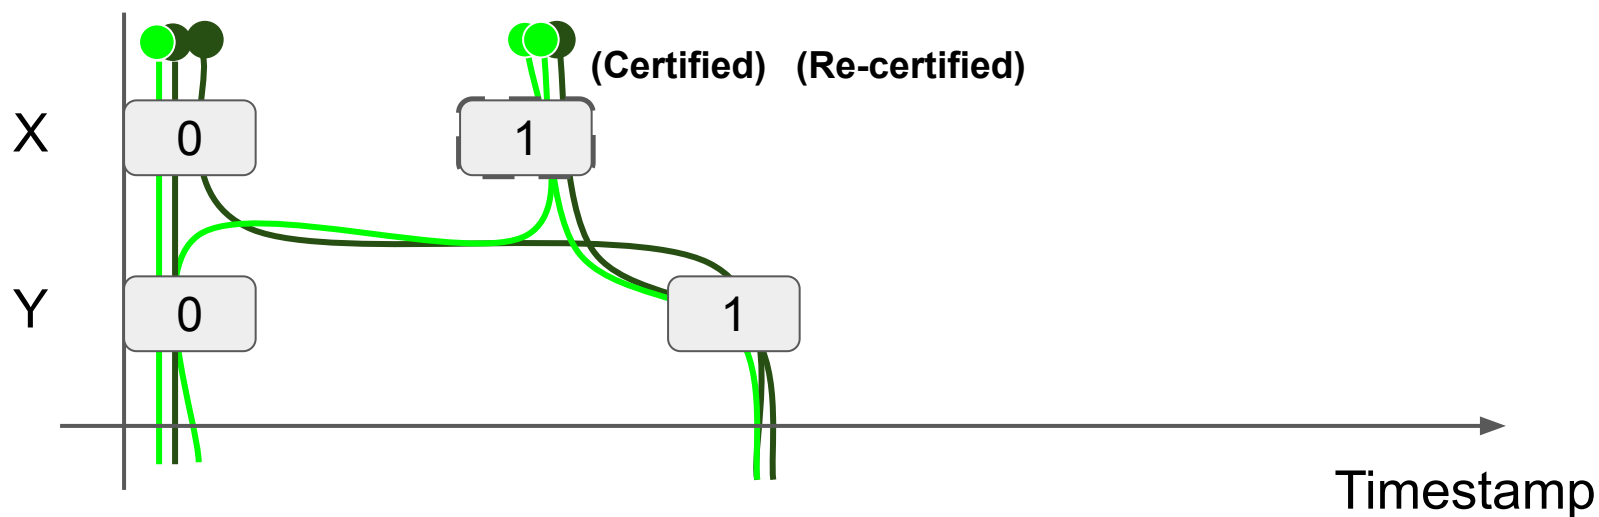
# Promising semantics 4: promises (3/5)

- **Store hoisting w/o dependency** [ $r1=r2=1$  **allowed** by reordering in the right]

○ ●  $r1=X$  || ●  $r2=Y$

●  $Y=r1$  || ●  $X=1$

● ●

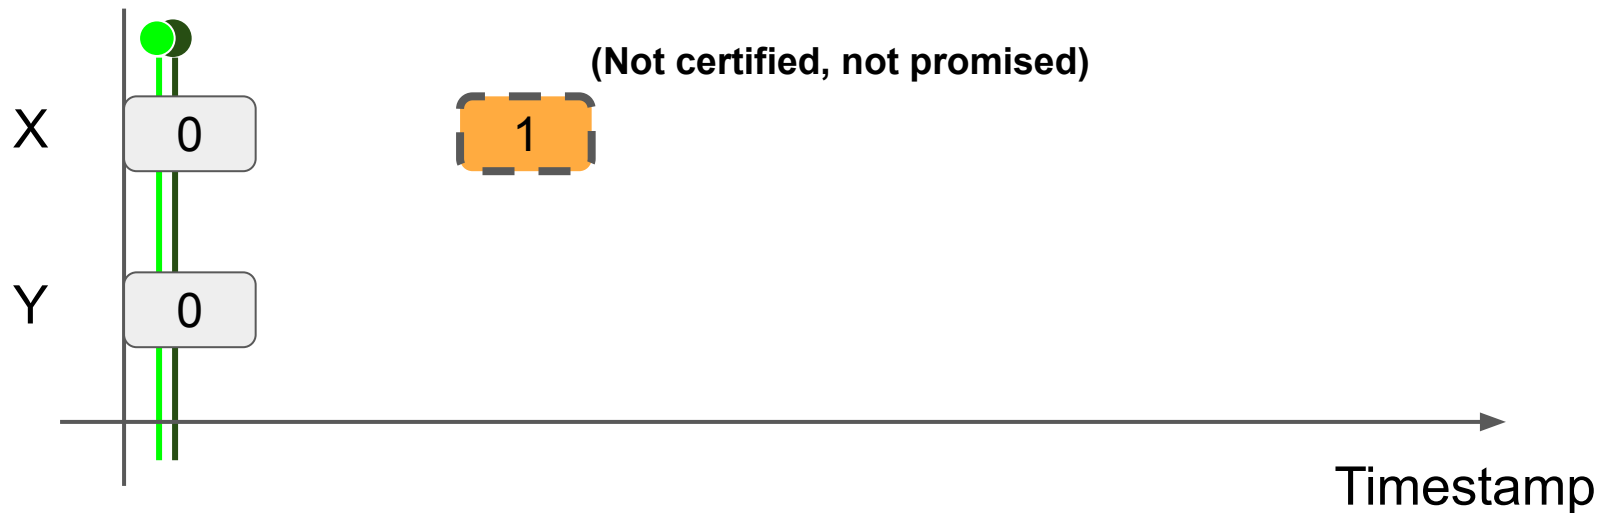


# Promising semantics 4: promises (4/5)

- Store hoisting w/ dependency [ $r1=r2=1$  **disallowed**, “out of thin air” (OOTA)]


$r1=X \quad || \quad r2=Y$

$Y=r1 \quad || \quad X=r2$



# Promising semantics 4: promises (5/5)



- **Store hoisting w/ syntactic dependency** [ $r1=r2=1$  **allowed**]
  - ●  $r1=X$  || ●  $r2=Y$   
 $Y=r1$  ||  $\text{if } r2==1 \{ X=r2 \}$   
 $\text{else } \{ X=1 \}$





# Promising semantics 4: promises (5/5), certification

- Store hoisting w/ syntactic dependency [ $r1=r2=1$  allowed]

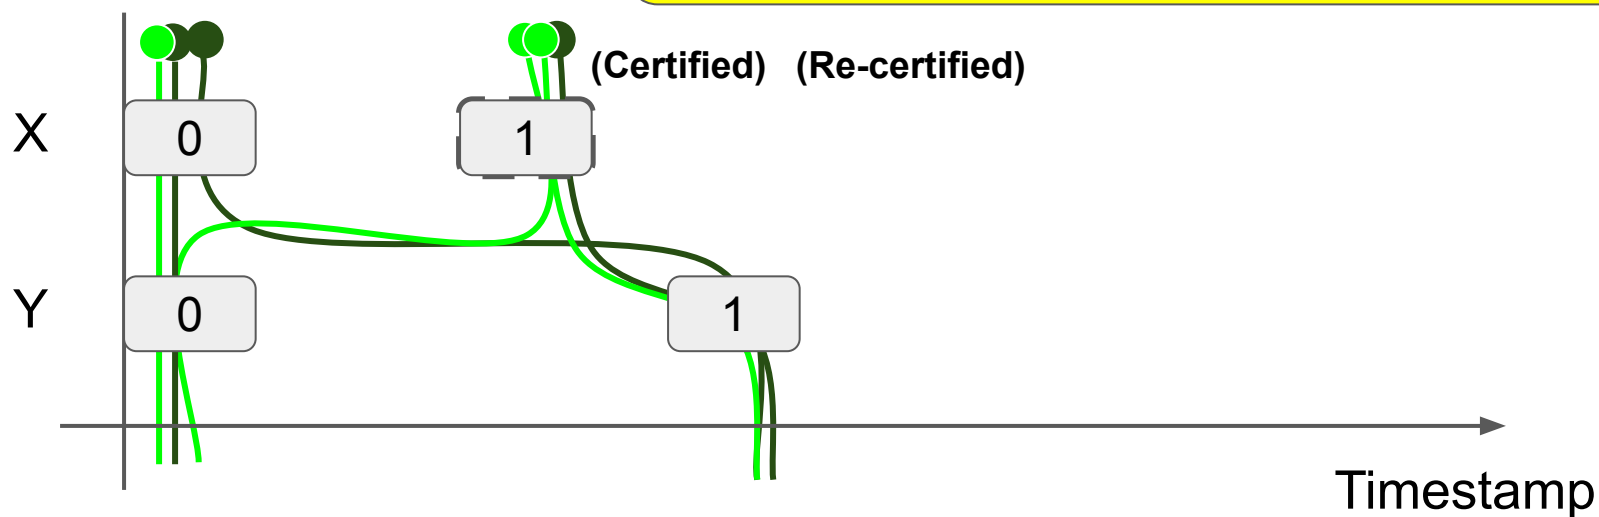


# Promising semantics 4: promises (5/5)

- Store hoisting w/ syntactic dependency [ $r1=r2=1$  allowed]

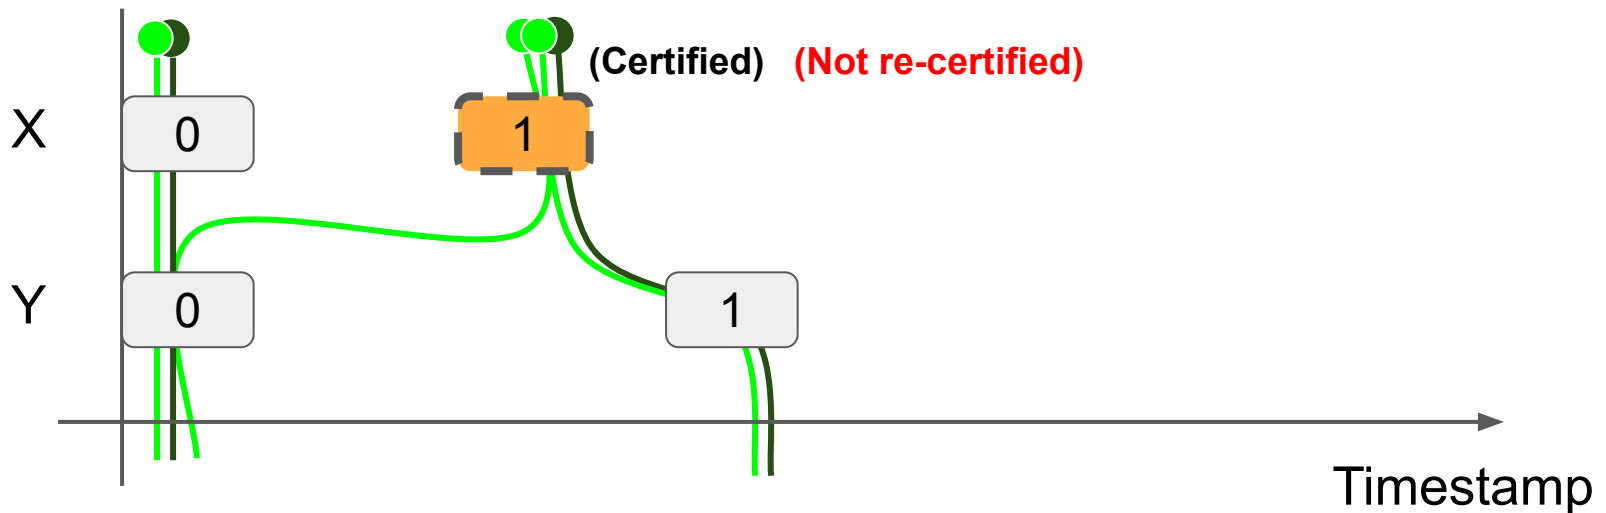
○ ●  $r1=X$  || ●  $r2=Y$   
●  $Y=r1$  || ●  $\text{if } r2==1 \{ X=r2 \}$   
●  $\text{else } \{ X=1 \}$   
●

**Certification and real execution may diverge to capture semantic dependency only**



# Bonus: behaviors constrained by re-certification

- **Store hoisting w/ syntactic dependency** [ $r1=r2=r3=1$  **disallowed**]
  - ●  $r1=X$  || ●  $r2=Y$ ;  $r3=X$  // due to RW coherence
  - $Y=r1$  || ●  $\text{if } r2==1 \{ X=r2 \}$
  - $\text{else } \{ X=1 \}$



# Summary of promising semantics

- An operational semantics modeling relaxed behaviors and orderings
- **Key ideas**
  - **Multi-valued memory:** modeling load hoisting
  - **Message adjacency:** modeling read-modify-write
  - **Views:** modeling coherence and synchronization
  - **Promises:** modeling store hoisting (**covered** in the final exam)
- **Homework:** making sense of the examples in the promising semantics
  - Message passing
  - Load hoisting
  - Store hoisting w/ & w/o (syntactic) dependency
  - [Java causality test cases](#) (#16, #19, #20 are wrong) — [already done!](#)

# Bonus: more on promising semantics

- <https://sf.snu.ac.kr/promise-concurrency/> for more details
- Designated courses for details on lower-level (e.g. (micro-)architecture)
- Related work
  - View-based semantics for ARM/RISC-V architectures  
<https://sf.snu.ac.kr/promising-arm-riscv/>
  - View-based semantics for persistent memory  
<https://cp.kaist.ac.kr/pmem>
  - A program logic (formal reasoning principle) for promising semantics  
<https://people.mpi-sws.org/~viktor/papers/esop2018-slr.pdf>
- Research ideas
  - View-based semantics for interrupts, non-cacheables, I/O, PCIe, CXL, ...
  - Verified compilation for promising semantics
  - ...

# Relaxed-memory behaviors


## Part 2: revisiting locks and objects

# Spinlock Correctness, Revisited

- ```
pub fn lock(&self) { while self.inner.cas(false, true, acquire).is_err() {} }
```

```
pub fn unlock(&self) { self.inner.store(false, release); }
```
- If a lock has already been acquired, lock() will spin
- Events between lock & unlock are transferred via release/acquire synchronisation.
- When holding the lock, you'll access the latest value of D (no shared access)



# Treiber's Stack Correctness, Revisited

- **Invariant:** head's pointer value has release view  $\sqsubseteq$  ( $\geq$  for each loc.)  
the messages of all node's value & next pointer
- **Release:** for maintaining the invariant (L54)
- **Acquire:** for exploiting the invariant (L68; reads at L70,84 are safe)
- **L74 can be relaxed:** the messages of later node's value & next pointer  
are already released
  - **“Release sequence”:** a msg's relview is transferred to the adjacent msg
  - E.g. B's next pointer and data is released at Head
    - [Head  $\rightarrow$  B  $\rightarrow \perp$ ] //  $\therefore$  release at L50
    - $\Rightarrow$  [Head  $\rightarrow$  A  $\rightarrow$  B  $\rightarrow \perp$ ] //  $\therefore$  head is CASed
    - $\Rightarrow$  [Head  $\rightarrow$  B  $\rightarrow \perp$ ] //  $\therefore$  head is CASed
    - $\Rightarrow$  [Head  $\rightarrow \perp$ ] //  $\therefore$  safe to read B's value and next pointer



# MS Queue Correctness, Revisited

- **Invariant:** a pointer value has release view  $\sqsubseteq$   
the messages of the pointed node's value & next pointer
  - A pointer value can be head, tail, or a node's next pointer
  - E.g. ptr's message has release view  $\sqsubseteq$  A's value and & pointer (to B)



- **Release:** for maintaining the invariant
  - MS queue: L88, 101, 111, 140, 148
- **Acquire:** for exploiting the invariant
  - MS queue: L77, 81, 125, 128)

# Sequence lock's synchronization

- W(0) happens before R(2): release/acquire synchronization
- R(2) doesn't see W(2)'s writes: release/acquire **fence** synchronization

|                      |        |                      |
|----------------------|--------|----------------------|
| // R(2)'s end        |        | // W(2)'s beginning  |
| ... // reading value |        | update seq 2->3      |
| fence(acquire)       | ←    → | fence(release)       |
| read seq (== 2?)     |        | ... // writing value |

(**Red**) if R(2) reads from any of W(2)'s writes,

(**Blue**) R(2)'s acquire fence happens after W(2)'s release fence,

(**Green**) invalidating R(2).

- Reads of R(2) and writes of W(2) should be “atomic” (no undefined behavior)
- R(2) may observe inconsistent state partially modified by W(2)  
(R(2) needs to be **resilient to such inconsistency**)

# Closing remarks

# What did you learn?

- **Concurrency:** about shared mutable resources
  - Difficult due to interleaving's nondeterminism
- **Shared memory:** the most widely-used shared mutable resources
  - More difficult due to reordering's nondeterminism
- **Design patterns**, both low-level (e.g. rel/acq) and high-level (e.g. helping)
- **Implementation** of locks, concurrent data structures, garbage collectors
  - Motivation, use cases, applications, tradeoffs, pitfalls
- **Promising semantics** for shared-memory concurrency

# Expected outcome achieved?

- **Don't be scared of concurrent programming**
  - In theory, it's just dealing with shared mutable states
- **Don't be scared of systems programming**
  - Systems programming: about low-level resources
  - Especially for parallel systems
- **Don't be scared of programming**
  - You'll learn systematic approach to concurrent programming
  - Which is also beneficial for programming in general

# What can I learn more about concurrency?

- **Hardware concurrency**

- Interrupt, I/O, system calls, persistent memory, ...
- X86, ARM, RISC-V concurrency
- Hardware description w/ Verilog, VHDL, [ShakeFlow](#)
- Compiler correctness

- **Concurrent data structures**

- ... (cannot summarize)
- OS kernel, DBMS, ...

Visit <https://cp.kaist.ac.kr>  
Contact [jeehoon.kang@kaist.ac.kr](mailto:jeehoon.kang@kaist.ac.kr)

- **Concurrent separation logic for verification**

- “Separation”: dealing with exclusive ownership of a part of resources
- The most widely-used reasoning principle for concurrency
- Iris: the most advanced CSL (<https://iris-project.org/>)

Thank you

# **FINAL EXAM**

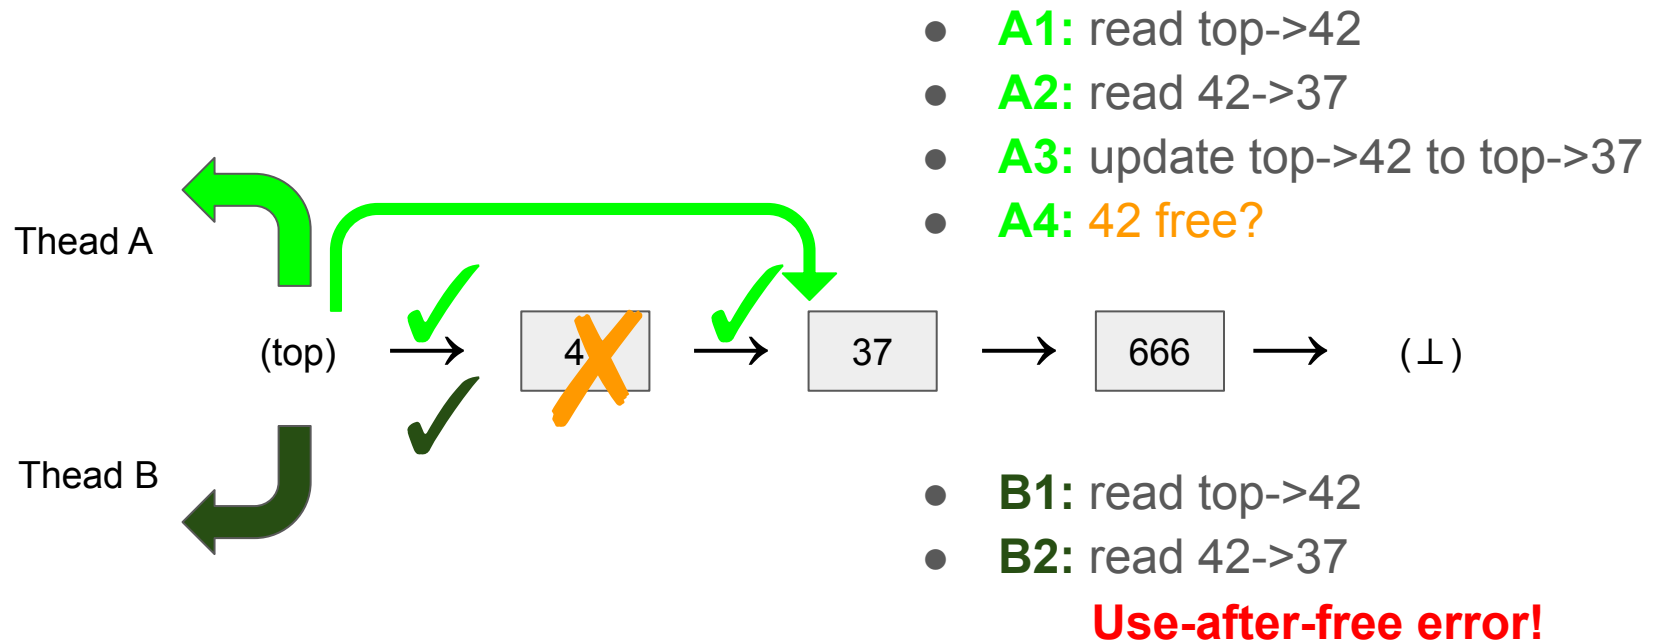
**Everything you've learned after midterm**



# Backup slides

# Example: memory reclamation in Treiber's stack

- Treiber's stack: singly linked list w/ head = stack top



**Question: how to ensure 42 is freed after finished being accessed?**


# Answer: by synchronization among threads

- **Memory reclaimer:** library/runtime dedicated for such synchronization
  - Motivation: too costly to solve the problem for each data structure
- **Many reclamation schemes**
  - Pointer-based reclamation (PBR)
  - Epoch-based reclamation (EBR)
  - Hybrids: QSBR, Snowflake, QSense, IBR, Hazard Era, **PEBR**, ...
- **Key ideas**
  - **Protect** blocks being accessed (by each thread)
  - **Retire** unlinked blocks (instead of immediately reclaiming them)
  - **Reclaim** retired blocks that are not protected by any threads

# Hazard pointer (HP)

- **Key idea 1: protect** blocks before accessing them
- **Key idea 2:** retire blocks  
Instead of immediately reclaiming them
- Reclaim those retired blocks that are not protected by any threads

# HP's example: Treiber's stack



- **A1:** ... (detaching 42)
- **A2:** retire 42
- **A3:** reclaim unprotected blocks

- **B1:** read top->42
- **B2:** protect 42를 보호
- **B3:** validate top->42
- **B4:** read 42->37






- **Case 1:** B2 => A2
  - B2 => A2 => A3
  - A3 doesn't reclaim 42

- **Case 2:** A2 => B2
  - A1 => A2 => B2 => B3
  - B3 cannot read top->42
  - B4 doesn't read 42

# HP's API

- **Data**
  - PL: Per-thread protected block (hazard pointer) list
  - RL: Per-thread retired block list
- **Protect(block)**
  - P1: add block to PL
  - P2: SC fence
  - P3: validate if the block is still pointed to by the pointer; otherwise, retry
- **Retire(block)**
  - R1: remove all references of the block from the shared memory
  - R2: SC fence
  - R3: add block to RL
- **Collect()**
  - C1: read every thread's PL
  - C2: reclaim those blocks in RL that are not in every thread's PL

# HP's problem: expensive synchronization


- **Synchronization points:** Inserting to RL , Inserting to PL 
  - To order  and 
  - Reordering may happen without synchronization
- **Synchronization cost:** fence in both A and B
  - 100+ cycles (e.g.: x86 mfence, ARMv8 dmb sy, POWER sync)
  - **Critical:** each traversal executes a fence in  ( $\frac{1}{3}$  throughput)

**Epoch-based reclamation: reducing synchronization cost**

# Epoch-based reclamation (EBR)



- **Key idea 1: critical region** of multiple accesses




- **Key idea 2: epoch consensus** (concurrent critical regions have “similar” epochs)
- **Protect epochs**, not pointers (amortizing synch. cost)



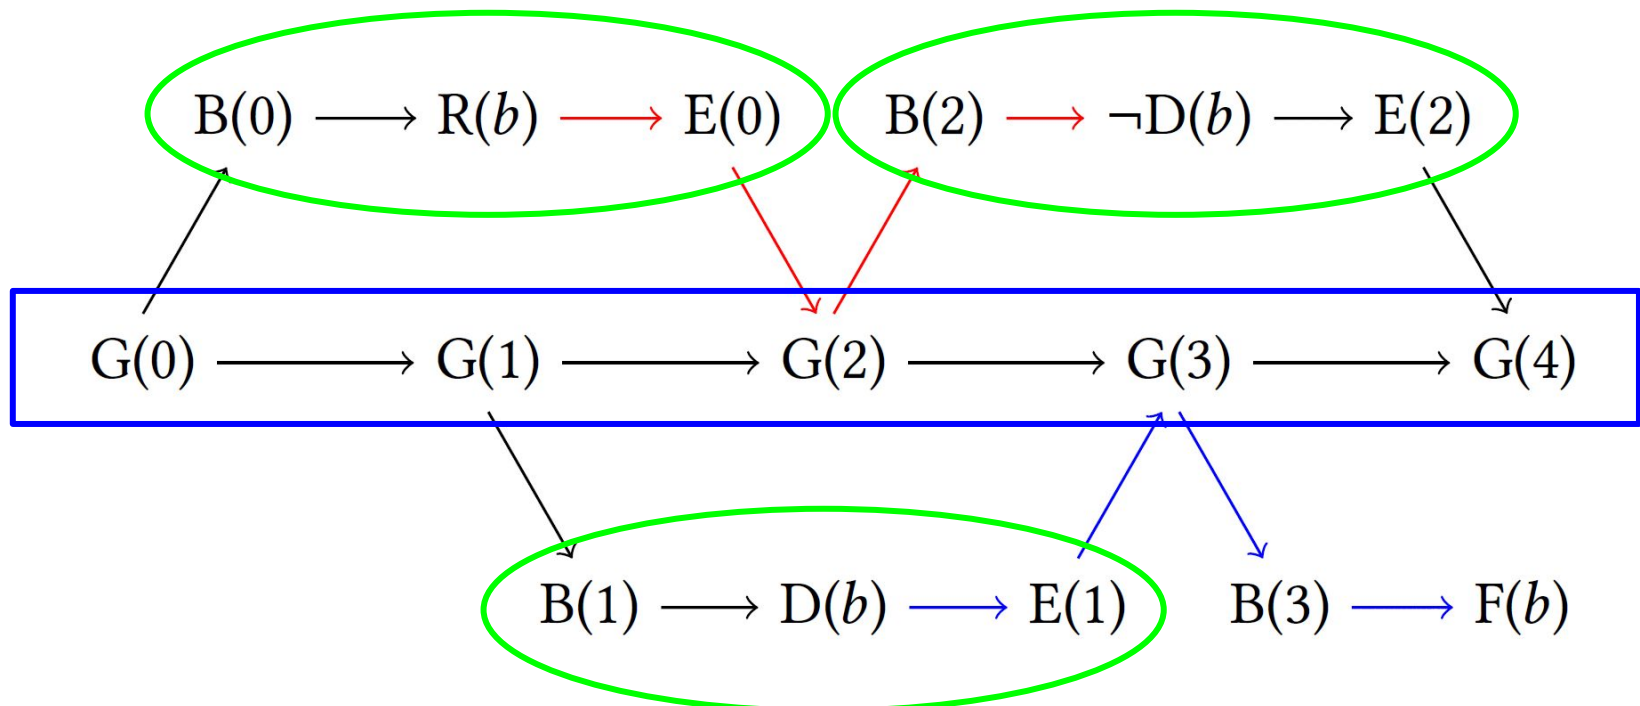
# EBR's key idea 1: critical region

- **Critical region (c.r.):** memory-accessing period inside a thread
  - **User-defined:** may differ among data structures and applications
    - E.g. during "stack.pop()", during processing a DB query
  - **Memory traversal within a critical region**
    - Shared memory access should be inside a c.r.
    - A pointer read inside a c.r. may be dereference inside the same c.r.
    - E.g. stack.pop() should be inside a c.r.
- **A1:** read top->42
  - **A2:** read 42->37
  - **A3:** update top->42 to top->37





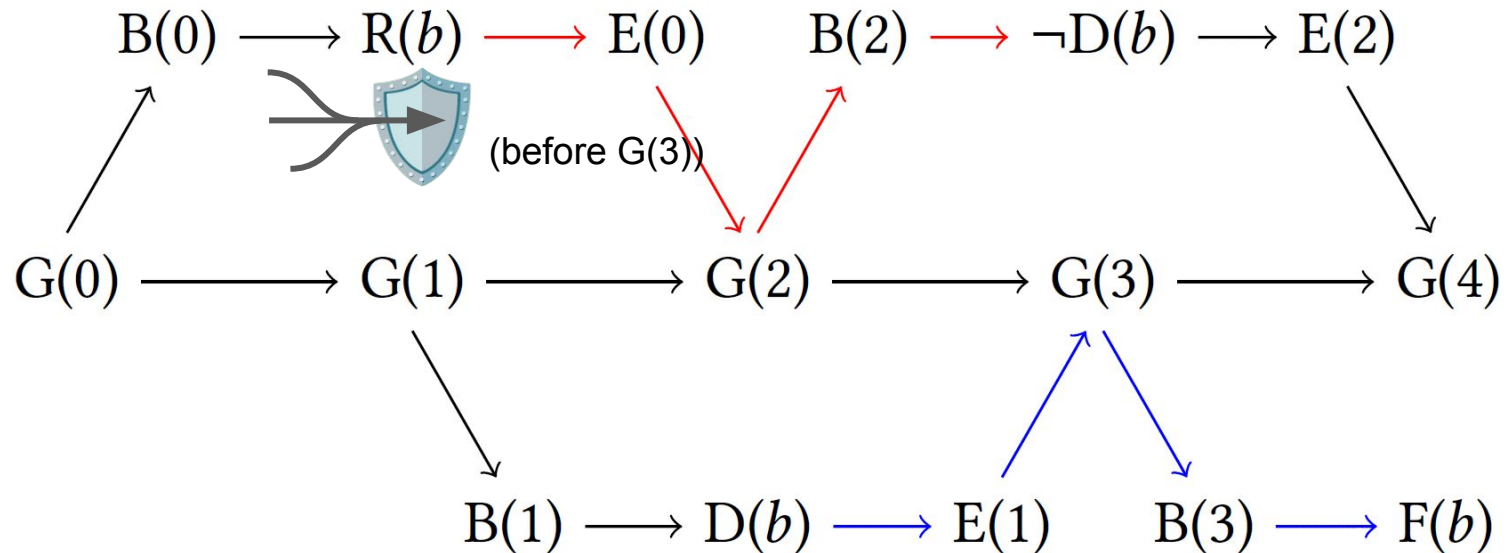
## EBR's key idea 2: epoch (1/2)

- **Epoch:** assigned to each c.r. (0, 1, 2, 3, ...),  $\rightarrow$ : happens-before
- **B(i), E(i):** the beginning/end of c.r. with epoch i, **G(i):** the beginning of epoch i
- **Epoch consensus:**  $G(i) \rightarrow B(i)$ ,  $E(i) \rightarrow G(i+2)$ 
  - Corollary: epochs of concurrent critical regions may differ only by 1
- **SC fence:** only at the beginning of critical regions



## EBR's key idea 2: epoch (2/2)

- **Protect:** blocks retired  at  $i$  are protected until  $G(i+3)$  
- **Safety:** (1) **inaccessible** in c.r. @  $i+2$ , (2) **live** in c.r. @  $i+1$
- **R(b):** retirement of  $b$  after detaching it from memory  
**D(b):** dereference of  $b$   
**F(b):** free(reclamation) of  $b$



# EBR's problem: not robust (memory leak)

- **When:** long critical region hinders epoch advancement and reclamation
  - Because  $E(i) \rightarrow G(i+2)$
- **Case 1:** user-defined long c.r. (e.g. OLAP, object cache, I/O)
- **Case 2:** unscheduled threads (e.g. oversubscription)
- **Case 3:** stalled threads (e.g. bugs, crash-only distributed systems)

**Question: fast & robust memory reclamation?**

# PEBR: marriage of HP and EBR

- **Goal:** fast like EBR and robust like HP
- **Key idea: hybrid of HP and EBR**
  - EBR (fast) at first, ejected to PBR (robust) when blocks are not reclaimed
- (2019) Jeehoon Kang and Jaehwang Jung.  
A marriage of pointer- and epoch-based reclamation. Submitted.  
<https://cp.kaist.ac.kr/gc/>

Shared memory  
low-level  
synchronization pattern

# Shared memory low-level synchronization patterns

- <https://jeehoonkang.github.io/2017/08/23/synchronization-patterns.html>
- **Pattern 1: “positive” release/acquire synchronization**
  - Positive: if A, then B
  - If a release store is read by an acquire load, view is transferred
- **Pattern 2: “contrapositive” release/acquire synchronization**
  - Contrapositive: if not B, then not A
  - If view is not transferred, a release store is not read by an acquire load
- **Pattern 3: “negative” SC synchronization**
  - Negative: either A or B
  - Either F1 happens before F2 or F2 happens before F1

**Covering “most” low-level synchronizations**

# Positive release/acquire synchronization

- If a release store is read by an acquire load,  
Then the view "released" at the store is "acquired" at the load
- E.g. message passing (X=1 should be acknowledged after reading Y=1)
  - ```
X = 1;                ||  if Y.load(acquire) {  
    Y.store(1, release); ||      assert(X == 1); // should not fail  
                        ||  }
```
  - Transferring data (X) w/ release/acquire synchronization of flag (Y)
- **The most widely-used low-level synchronization pattern**
  - Used in spinlock, channel, ...



# Negative SC synchronization

- All SC fences are strictly ordered w.r.t. the per-thread views
- E.g. message passing (X=1 should be acknowledged after reading Y=1)
  - ```
X = 1;                ||  if Y.load(relaxed) {  
fence(SC);           ||      fence(SC);  
Y.store(1, relaxed); ||      assert(X == 1); }
```
  - Transferring data (X) or ignoring flag (Y) w/ SC synchronization
- **An advanced low-level synchronization pattern**
  - Used in Peterson's mutex, memory reclamator, work-stealing deque, ...

# Example: Peterson's mutex (implementation)

- `let flag: [AtomicBool; 2]; // whether a thread wants to begin a critical section`  
`let turn: AtomicUsize = 0; // who has the precedence?`
- `fn begin(id: Usize) { // thread id: 0 or 1 (two threads, T0 and T1)`  
    `flag[id].store(true);`  
    `fence(SeqCst); // A`  
    `turn.store(1 - id);`  
    `fence(SeqCst); // B`  
    `while (flag[1 - id].load(acquire) && turn.load() == 1 - id) {}`  
}
- `fn end(id: Usize) {`  
    `flag[id].store(false, release);`  
}

# Example: Peterson's mutex (correctness)

- Case analysis for the order of fences
- Case 1: **A0 -> B0 -> A1 -> B1**
  - **flag[0] = true** and **turn = 1** should be ack. at A1
  - **turn = 1** before **turn = 0** w.r.t. Coherence
- Case 2: **A0 -> A1 -> B0 -> B1** or **A0 -> A1 -> B1 -> B0**
  - Both **flag[0] = true** and **flag[1] = true** should be ack. at **B0** and **B1**
  - W/o loss of generality, assume **turn = 1** before **turn = 0** w.r.t coherence
- ... (similar cases)
- **flag[0] = true** is ack. at B1 && **turn = 1** before **turn = 0** w.r.t coherence
- T1 should spin until **T0** calls end() writing **flag[0] = false**
- Mutex thanks to release/acquire synch. from **T0** to T1 via flag[0]
- Homework: Is each of the fences (A and B) necessary?

# Example: Peterson's mutex (discussion)

- **Not practical due to drawbacks**
  - Unnecessarily complicated
  - Not reusable: cannot begin() after end() (hence not “lock”)
- **Theorem:** reusable mutex requires atomic rmw instructions  
(e.g. swap, compare-and-swap, or other rmw instructions)
  - Peterson's mutex is an artifact before rmw instructions
- **Next section: more lock implementations w/ tradeoffs**
  - Spinlock's advantages: simple, fast (when uncontended), compact
  - Spinlock's drawbacks: unfair, inscalable, energy-inefficient

# Videos on data parallelism and async I/O

# Rayon: Data Parallelism for Fun and Profit

- Nicholas Matsakis @ Rust Belt Rust 2016
- (Turn on English caption)



The slide displays a transformation of Rust code to utilize Rayon for parallelism. The top code block shows a sequential function `increment_all` that iterates over a mutable array `counts` and increments each element. A blue arrow points down to a second code block where the iteration is replaced by `paths.par_iter_mut().for_each(|c| *c += 1);`. A red text annotation states: `'c' not shared between iterations!`. In the bottom right corner, there is a small video inset showing a man gesturing while speaking.

```
fn increment_all(counts: &mut [u32]) {  
    for c in counts.iter_mut() {  
        *c += 1;  
    }  
}
```

↓

```
fn increment_all(counts: &mut [u32]) {  
    paths.par_iter_mut()  
        .for_each(|c| *c += 1);  
}
```

`'c' not shared  
between iterations!`

# Zero-Cost Async IO

- Boats @ Rust LATAM 2019
- (Turn on English caption)




# Optimistic lock coupling



## Lock coupling (“hand-over-hand locking”)


- **Holding at most 2 locks at a time during traversal**
  - Each node has its own lock
  - During traversal, hold locks for the “current” node and its parent
  - Holding the lock for the parent to protect structural changes



## Read operations invalidates cache!

# Optimistic lock coupling (OLC)

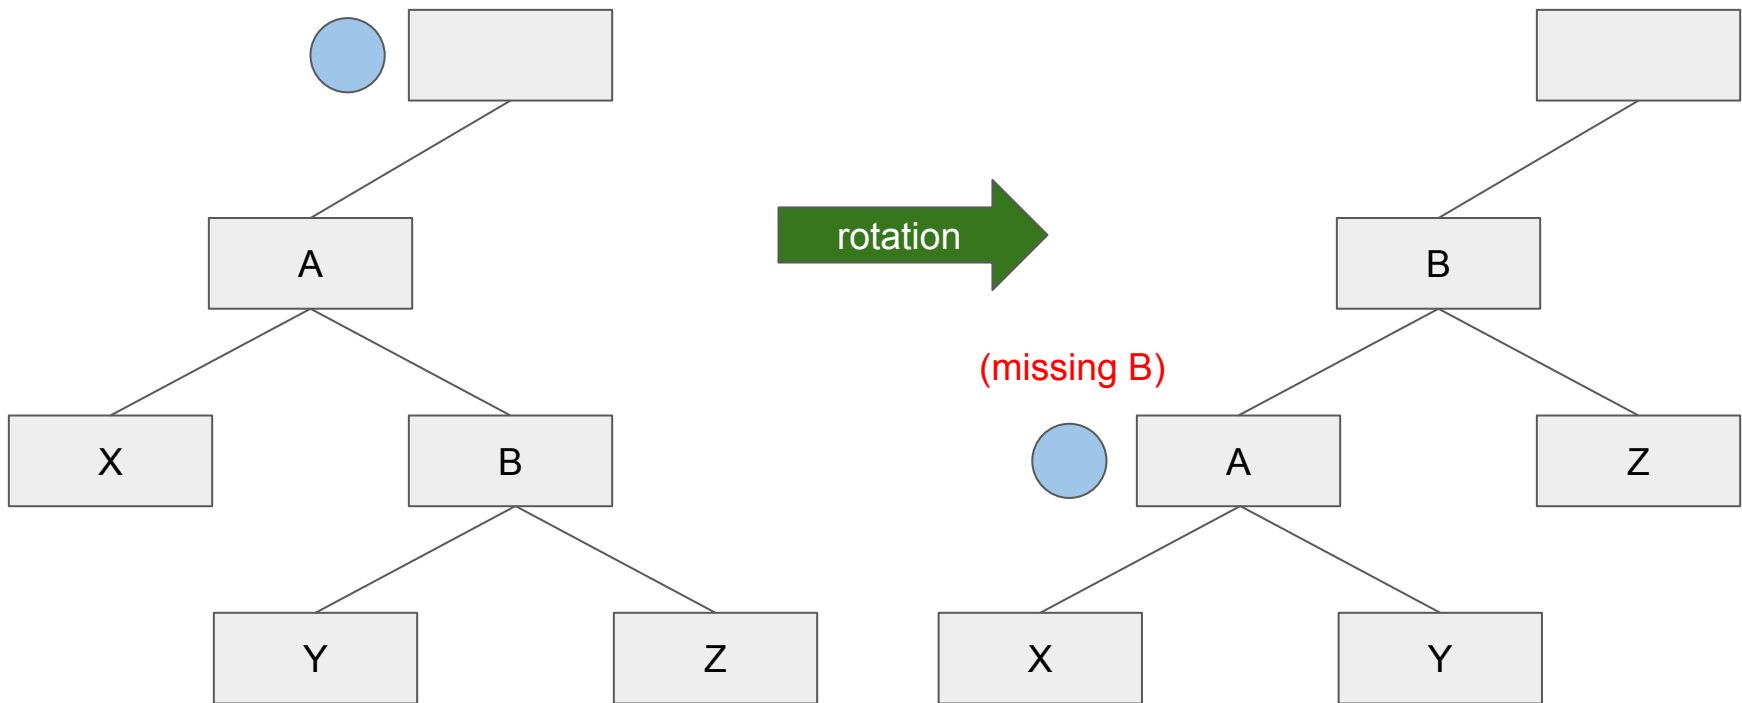
- **Holding at most 2 sequence locks at a time during traversal**
  - Upgrading to write lock to protect structural changes



**Read operations  
do not invalidate cache!**

# Why acquiring locks of two nodes simultaneously?

- Otherwise, we may miss values (e.g. B below) during traversal.



# When is it safe to deallocate a node?

- **Question:** Read doesn't leave any footprint on memory. How to make sure concurrent threads don't deallocate the accessed node?
- **Approach:** separate synchronization mechanism for safe deallocation
  - “Safe memory reclamation scheme” or just garbage collection
  - E.g. epoch-based reclamation, hazard pointers, ...
  - Will be discussed later in this course
- **Let's assume epoch-based reclamation (EBR)**
  - Operations are delimited by “critical sections”
  - Memory blocks are accessed and *retired* only inside critical sections
  - Retired blocks are automatically deallocated **later when safe**
  - Implementation: [Crossbeam](#) (concurrency library in Rust)

# Old Annoucements

# Homework 1 (due: Nov. 4, 20% of the credits)

- **Implement (sequential) adaptive radix tree in Rust**
  - <https://github.com/kaist-cp/cs431/issues/6>
  - Read paper: <https://db.in.tum.de/~leis/papers/ART.pdf>
  - Read skeleton: [\[link\]](#)
- **Read skeleton's README.md for specifications and recommendations**

# Homework 2 (due: Nov. 29, 20% of the credits)

- **Synchronize binary search tree with OLC**
  - **IMPORTANT:** not AVL, but plain binary search tree
  - Read paper: <https://stanford-ppl.github.io/website/papers/ppopp207-bronson.pdf>
  - Read skeleton: [\[link\]](#)
  - Grading: TBA
- **Tip: read paper early!**

# Midterm exam (Oct. 21, 20% of the credits)

- **Date & time:** 9:00am-11:00am, October 21th, 2019
- **Place:** Rm. 2111, Bldg. E3-1
- **Coverage:** everything you've learned in this course
- **Bring a black/blue pen and your student ID**
- **Study guide**
  - You will be asked to “explain” the reasons of a phenomena. Your explanation should be an informal proof, but with a few gaps. Your informal proof should be clear and precise.
  - You will be asked about the implementations details of the code presented in class (e.g. locks and ART). You may be asked to hand-write an implementation.
  - You will be asked about the video lecture.
- **Questions or comments?**



# Final exam (Dec. 16, 20% of the credits)

- **Date & time:** 9:00am-11:00am, December 16th, 2019
- **Place:** Rm. 2111, Bldg. E3-1
- **Coverage:** everything you've learned after the midterm exam
- **Bring a black/blue pen and your student ID**
- **Study guide**
  - You will be asked to “explain” the reasons of a phenomena. Your explanation should be an informal proof, but with a few gaps. Your informal proof should be clear and precise.
  - You will be asked about the implementations details of the code presented in class. You may be asked to hand-write an implementation.
  - You will be asked about the video lecture.
- **Questions or comments?**

# Garbage slides

# Nondeterminism due to thread interleaving (TODO)

- **E.g.** concurrent counter: multiple threads incrementing a shared location:

```
static COUNTER = AtomicUsize::new(0);  
// thread A & B  
let c = COUNTER.load();  
COUNTER.store(c + 1);
```
- **Problem:** unintended behaviors due to unfortunate scheduling
  - [COUNTER=0] A load, B load, A store, B store [COUNTER=1]
- **Solution:** forbidding such a scheduling by **atomically reading & writing**
  - “**Read-modify-write**”, e.g. swap, compare-and-swap, fetch-and-add
  - // thread A & B

```
let c = COUNTER.fetch_and_add(1);
```
  - [COUNTER=0] A fetch\_and\_add, B fetch\_and\_add [COUNTER=2]

# Optimizations introducing? relaxed behaviors



- **Reordering:** unless accessing the same location,  
    **any** two load/store/rmw instructions can be reordered.
  - E.g. `X=1; r=Y -> r=Y; X=1` (load-store reordering)
  - E.g. `X=1; Y=1 -> Y=1; X=1` (store-store reordering)
- **Merging:** if accessing the same location, TODO: not introducing??  
    two load/store/rmw instructions **may** be merged.  
    (only when it makes sense for sequential program)
  - E.g. `X=1; X=2; r=X -> X=2; r=X` (store-store, store-load merging)
  - E.g. `X=1; X.fetch_and_add(1) -> X=2` (store-rmw merging)
- **Eliminating:** if the result of a load is not used, the load may be eliminated.
  - E.g. `r=X; -> nop` (load elimination) TODO: not introducing??

# Wait-free data structures

- Key idea: each thread publishes its ongoing operation, the “winning” thread helps the others by finishing the published operations by the others
- Helping: <https://dl.acm.org/citation.cfm?id=102808>
- Flat combining: <https://www.cs.bgu.ac.il/~hendlerd/papers/flat-combining.pdf>
- Example: [A Wait-free Queue as Fast as Fetch-and-Add](#)

# 포인터/시대 기반 수집기법 (PEBR)

- 아이디어: PBR과 EBR의 혼종



- 처음에는 EBR (빠름), 메모리 수집이 잘 안되면 PBR로 **강퇴** (수집 보장)

# Intermission: Crossbeam (TODO)

- [Crossbeam](#): most widely used concurrency library for Rust
- **Utilities** (e.g. CachePadded)
- **Pointer value manipulation APIs**
  - Owned: TODO
  - Atomic: TODO
  - Shared: TODO
- **Garbage collection APIs** (epoch-based reclamation)
  - Guard: TODO
  - `defer_destroy()`: TODO
- **lock-free data structures**
  - Stack, queue, work-stealing deque, channel, ...



# Pointer/epoch-based reclamation (PEBR)

- <https://cp.kaist.ac.kr/gc/>
- **Idea: hybrid of PBR and EBR**
  - EBR (fast) at first, ejected to PBR (robust) when blocks are not reclaimed
- **Results**
  - As fast as EBR (85%-90% throughput)
  - Robust as PBR (guaranteeing reclamation of blocks)
  - Portable, generally applicable to many data structures, compact
  - **The first scheme that satisfies all the criteria above (i.e. superior)**
- We will learn PBR and EBR (as background), and PEBR



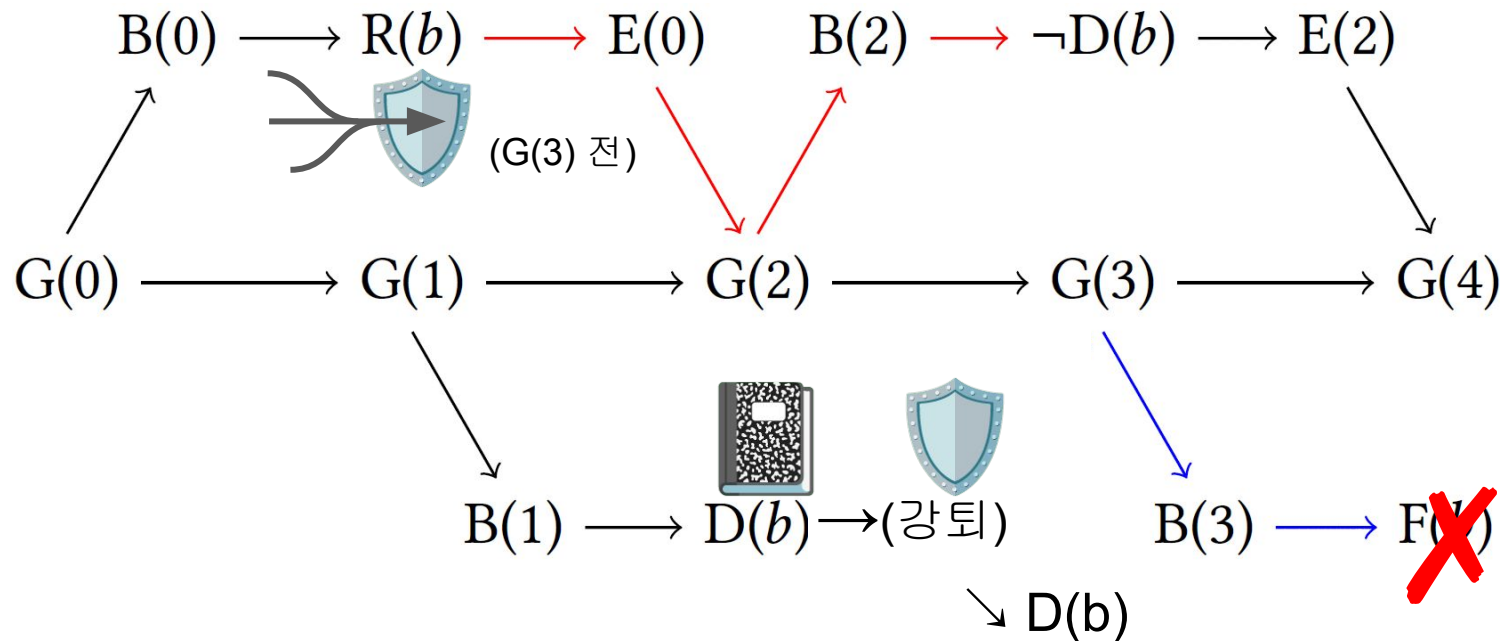
# 포인터/시대 기반 수집기법 (PEBR)

- 아이디어: PBR과 EBR의 혼종




- 처음에는 EBR (빠름), 메모리 수집이 잘 안되면 PBR로 **강퇴** (수집 보장)

# 포인터/시대 기반 수집기법 아이디어: 강퇴







- EBR에 기반, 묶음이 너무 길어지면 강퇴하고 새 시대 시작
- 강퇴시 묶음에서 접근한 포인터 다시 접근 가능,  $G(3)$  이후에도 해제하면 안됨
- 접근한 포인터를 (쓰레드 내에) 기록해뒀다가 강퇴시 보호

# 포인터/시대 기반 수집기법 장단점

- **장점 1:** 빠르고 메모리 수집 보장 (강퇴를 통해)
- **장점 2:** 강퇴 과정이 락없고 이식성 높음
- **장점 3:** 포인터를 오래 들고있어야 하는 경우 고의강퇴 가능
- **단점 1:** 접근한 포인터 기록  하는 비용 (10-15%)
- **단점 2:** 프로그래밍하기 PBR/EBR보다 약간 복잡함

**장점 1이 모든 단점을 커버하고도 남음**


# 포인터/시대 기반 수집기법 결과 1: Throughput




EBR이 메모리 수집 안해서

# 포인터/시대 기 Throughput이 적어 경향성 파악 실패 메모리 사용량

List




Hashmap




EBR, stalled는 메모리수집 아예 안함

EBR은 스레드 많으면 수집 느려짐

NMtree



Bonsai



SMR, interf.

- NR
- EBR
- EBR, 10ms
- EBR, stalled
- PEBR
- PEBR, 10ms
- PEBR, stalled

# Synchronization in HP: use-after-free never happens

- Suppose T1 retires and reclaims a block B, and T2 protects and uses B.
- **Case 1: T2's P2 happens before T1's R2**
  - P1's write to PL is visible to C2
  - To reclaim B, T1's C2 should acknowledge the removal of B from PL (T2 finishes accessing B and then remove B from PL)
  - T2's access to B happens before T1's reclamation of B
- **Case 2: T1's R2 happens before T2's P2**
  - P3's validation fails because it cannot read B from the memory
  - T2 cannot access B

# Analyzing shared mutable states with ownership

- **Application to concurrency:** fit for analyzing synchronization among threads
  - E.g., if thread A owns a buffer, thread B cannot access it
- **Its ownership type is an extremely useful abstraction for SMS**
  - **Statically proving the safety** of accesses to SMS
  - Not only for sequential but also for concurrent programs
- **Key idea 1: disallowing shared mutable operations (SMOs) by default**
  - A resource is either exclusive or immutable (but not both)
  - **Exclusive:** resource is read/written by its designated **owner**  
or its exclusive **borrower**; or
  - **Immutable:** resource is read by its shared borrowers
- **Key idea 2: allowing SMOs in a controlled way**
  - Motivation: SMOs are essential and unavoidable (e.g. in concurrency)
  - Mechanism: **interior mutability** via **unsafe** blocks

# Rust and concurrent programming

- Implementation w/ SMOs
  - Essential in concurrent programming
- Interface virtually w/o SMOs
  - For the ease of concurrent programming
- E.g. concurrent stack w/ non-SMO interface
  - `fn Stack::push(&self, value: T);`
  - `fn Stack::pop(&self) -> option<T>;`
- Using compile-time lifetime as a static verification of that of the critical sections (more later)



# Bonus: foundation of Rust's ownership type

- **Layer 3: ownership type**
  - Computational checker / mathematical “library” for proving the safety of accesses to SMS
- **Layer 2: lifetime logic**
  - Mathematical library for proving the safety specialized for lifetimes and borrows
- **Layer 1: concurrent separation logic (CSL)**
  - Mathematical proof system for proving the safety
  - Ownership type checker as a “lemma” for the safety proof in CSL
  - We'll see more later...

# Q: why are you so narcissistic...?

- This course deals with work by myself only.
- **Reason 1: it's the material I can teach at my best**
  - I know the in and out of the work, including pitfalls and key points
- **Reason 2: I want to deliver a coherent narrative about concurrency**
  - No existing such a coherent narrative
  - I needed to build up my own self (semantics, design patterns, library, ...)
  - This course: a first attempt